HFS Utilities

Complete Technical Manual

HFS and HFS+ Filesystem Utilities for Unix Systems

hfsutils Project

December 18, 2025

Version 4.1.0A.2

Copyright and License

This manual documents the HFS Utilities (hfsutils) project, a collection of tools for creating, checking, mounting, and manipulating HFS and HFS+ filesystems on Unix-like systems including Linux, BSD, and macOS.

License: GNU General Public License v2 or later

Project: https://github.com/JotaRandom/hfsutils

Maintainer: hfsutils Project Contributors

Version: 4.1.0A.2

Contents

Chapter 1

Filesystem History and Evolution

1.1 Overview

This chapter traces the evolution of Apple's filesystems from the original Macintosh File System (MFS) through HFS, HFS+, and the modern APFS. Understanding this evolution provides context for design decisions and compatibility requirements.

1.2 Macintosh File System (MFS)

1.2.1 Introduction

MFS (Macintosh File System) was the original filesystem used on the Macintosh 128K and 512K (1984-1985). It was designed for floppy disks and small hard drives.

1.2.2 Key Characteristics

- **Introduced**: January 1984 (Macintosh 128K)

- **Maximum volume size**: 20 MB (limited by hardware)

- **Maximum files**: 65,535 (16-bit file numbers)

- **Filename length**: 255 characters (Macintosh Roman)

- **Directory structure**: Flat (no folders/subdirectories)

- **Dual forks**: Data fork and resource fork

- **File types**: 4-character type and creator codes

1.2.3 Design Philosophy

MFS was revolutionary for its time:

- First consumer filesystem with resource forks

- Integrated with Finder for desktop metaphor

- Type/creator codes enabled double-click launching

- Desktop file tracked file positions and icons

1.2.4 Limitations

- **No hierarchy**: All files stored in single directory

- **Finder folders**: Simulated by Finder, not filesystem

- **Small volumes**: Designed for 400K floppies

- **Poor scalability**: Performance degraded with many files

1.2.5 MFS Structure Summary

Component	Description
Logical Block 0-1	Boot blocks (1024 bytes)
Logical Block 2	Master Directory Block (MDB)
Following blocks	File directory (single B-tree)
...	Allocation bitmap
...	File data
Last sector	Alternate MDB

Table 1.1: MFS Volume Layout

Note: MFS is not implemented in hfsutils but documented for historical context and understanding HFS evolution.

1.3 Hierarchical File System (HFS)

1.3.1 Introduction

HFS (Hierarchical File System) replaced MFS in Mac OS System 2.1 (1985) and System 3.0 (1986). It introduced true hierarchical directories.

1.3.2 Key Improvements over MFS

- **Hierarchical directories**: True folder structure

- **Larger volumes**: Up to 2 GB (later 2 TB with HFS wrapper)

- **Better performance**: Separate B-trees for catalog and extents

- **Allocation optimization**: Clump sizes reduce fragmentation

1.3.3 Timeline

- **1985**: Introduced with Hard Disk 20

- **1986**: Became standard in System 3.0

- **1998**: Superseded by HFS+ in Mac OS 8.1

- **1999**: Last use in Mac OS 9 for booting

- **2000s**: Maintained for compatibility

- **Today**: Supported for legacy media

1.3.4 HFS Characteristics

Detailed in Chapter 2 (HFS Specification). Summary:

Feature	Specification
Signature	0x4244 ('BD')
Maximum volume	2 TB
Maximum file	2 GB
Filename length	31 characters (MacRoman)
Allocation blocks	16-bit addressing (65,535 blocks)
B-tree node size	512 bytes (fixed)
Date range	1904-2028 (Y2K28 limit)
Case sensitivity	Case-insensitive, case-preserving

Table 1.2: HFS Characteristics

1.4 HFS Plus (HFS+)

1.4.1 Introduction

HFS+ (HFS Plus, "Mac OS Extended") was introduced in Mac OS 8.1 (1998) to address HFS limitations for modern computing.

1.4.2 Major Enhancements

- **Unicode filenames**: UTF-16, up to 255 characters

- **32-bit addressing**: Support for very large volumes

- **Smaller allocation blocks**: Better space efficiency

- **Journaling**: Optional crash recovery (Mac OS X 10.2.2+)

- **Hard links**: Unix-style hard links

- **Symbolic links**: Unix-style symbolic links

- **Extended attributes**: Arbitrary metadata

- **HFSX variant**: Case-sensitive option

1.4.3 Timeline

- **1998**: Introduced in Mac OS 8.1

- **1999**: Became default in Mac OS 8.6

- **2001**: Mac OS X adoption

- **2002**: Journaling added (Mac OS X 10.2.2)

- **2005**: Case-sensitive HFSX variant

- **2017**: Superseded by APFS (macOS 10.13)

- **Today**: Still widely used, especially on spinning drives

1.4.4 HFS+ Variants

Standard HFS+

- Signature: 0x482B ('H+')

- Version: 4

- Case-insensitive filenames

- Most compatible

HFSX (HFS Extended)

- Signature: 0x4858 ('HX')

- Version: 5

- Case-sensitive filenames

- Used for Unix-like behavior

Journaled HFS+

- Attributes bit 13 (0x2000) set

- Journal info block pointer in Volume Header

- Circular journal buffer

- **Not supported by Linux kernel**

1.4.5 HFS+ Characteristics

Detailed in Chapter 3 (HFS+ Specification). Summary:

Feature	Specification
Signature	0x482B ('H+') or 0x4858 ('HX')
Maximum volume	8 EB theoretical
Maximum file	8 EB theoretical
Filename length	255 UTF-16 characters
Allocation blocks	32-bit addressing
B-tree node size	4096 bytes (typical)
Date range	1904-2040 (Y2K40 limit)
Case sensitivity	Optional (HFSX)
Journaling	Optional

Table 1.3: HFS+ Characteristics

1.5 Apple File System (APFS)

1.5.1 Introduction

APFS (Apple File System) is Apple's modern filesystem, introduced in macOS 10.13 High Sierra (2017). It replaces HFS+ as the default for SSDs and flash storage.

1.5.2 Design Goals

- **Flash optimized**: Minimize write amplification

- **Space sharing**: Multiple volumes share space pool

- **Snapshots**: Instant, space-efficient snapshots

- **Cloning**: Copy-on-write file/directory clones

- **Strong encryption**: Native full-disk and per-file encryption

- **Crash protection**: Copy-on-write metadata

1.5.3 Key Features

- 64-bit inode numbers

- Nanosecond timestamp precision

- Space sharing between volumes

- Native encryption (FileVault)

- Fast directory sizing

- Atomic safe-save operations

- Clones and snapshots

1.5.4 APFS vs HFS+

Feature	HFS+	APFS
Introduced	1998	2017
Optimized for	HDDs	SSDs/Flash
Snapshots	No	Yes
Cloning	No	Yes (instant)
Encryption	Per-volume	Per-file + volume
Space sharing	No	Yes
Timestamps	Second precision	Nanosecond precision
Max files	4 billion	9 quintillion
Journaling	Optional	Always (COW)

Table 1.4: HFS+ vs APFS Comparison

1.5.5 APFS Adoption

- **macOS 10.13+**: Default for SSDs

- **iOS 10.3+**: Default for all devices

- **watchOS 4+**: Default

- **tvOS 11+**: Default

- **HDDs**: HFS+ still recommended (as of macOS 13)

1.5.6 Why HFS+ Still Matters

- **Legacy systems**: Pre-2017 Macs

- **HDDs**: APFS not optimized for spinning drives

- **External drives**: Better compatibility

- **Recovery partitions**: Some use HFS+

- **Boot Camp**: Windows compatibility

- **Linux support**: Better HFS+ kernel support

1.6 Filesystem Evolution Summary

FS	Year	Max Volume	Max File	Key Innovation
MFS	1984	20 MB	N/A	Resource forks
HFS	1985	2 TB	2 GB	Hierarchy
HFS+	1998	8 EB	8 EB	Unicode, large files
APFS	2017	8 EB	8 EB	Flash-optimized, snapshots

Table 1.5: Apple Filesystem Evolution

1.7 Scope of This Document

This manual focuses on:

- **HFS (Classic)**: Complete implementation (Chapter 2)

- **HFS+**: Complete implementation (Chapter 3)

- **MFS**: Historical context only

- **APFS**: Overview for migration context

All specifications are documented at the bit level to enable complete reimplementation without external references.

1.8 Historical Oddities and Notes

1.8.1 The "BD" Signature Mystery

HFS uses signature 0x4244 ('BD' in ASCII). The origin is unclear:

- Possibly "Block Device"

- Possibly arbitrary choice

- Never officially documented by Apple

1.8.2 The 1904 Epoch

All Apple filesystems use January 1, 1904 as time zero:

- Predates Unix epoch (1970) by 66 years

- Reason unknown but consistent across all Apple filesystems

- Creates Y2K28 (HFS) and Y2K40 (HFS+) problems

1.8.3 The Colon as Path Separator

Classic Mac OS used colon (:) as path separator:

- Unix/Windows use / and \respectively

- Filenames cannot contain colons

- Causes issues when sharing files cross-platform

- Mac OS X translates : to / for display

1.8.4 The 31 vs 255 Character Limit

- HFS: 31 characters (historical, tied to original Mac)

- HFS+: 255 characters (modern standard)

- MFS: Actually 255 characters (more than HFS!)

1.8.5 Case Sensitivity Confusion

- HFS/HFS+: Case-insensitive by default

- HFSX: Case-sensitive variant

- Linux ext4: Case-sensitive

- Windows NTFS: Case-insensitive but preserving

- Source of many cross-platform file issues

Chapter 2

Introduction

2.1 Overview

The HFS Utilities (hfsutils) project provides a comprehensive set of tools for working with Apple's HFS (Hierarchical File System) and HFS+ (HFS Plus) filesystems on Unix-like operating systems including Linux, BSD, and macOS.

2.1.1 Purpose and Goals

HFS and HFS+ filesystems are commonly used on:

- Classic Mac OS systems (HFS)

- Mac OS X and macOS systems prior to APFS (HFS+)

- iPod devices

- External drives formatted for Mac compatibility

- Disk images and backups from Apple systems

This toolset enables Unix systems to:

- **Create** HFS and HFS+ filesystems with `mkfs.hfs` and `mkfs.hfs+`

- **Check and repair** filesystem integrity with `fsck.hfs` and `fsck.hfs+`

- **Mount** HFS and HFS+ volumes with `mount.hfs` and `mount.hfs+` (requires kernel support)

- **Manipulate files** on HFS volumes without mounting using `hfsutil` commands (useful on systems without HFS kernel drivers)

2.1.2 Supported Systems

The utilities work on any POSIX-compliant system with appropriate kernel support:
Note: On systems without kernel HFS support, the `hfsutil` commands can still be used to access HFS filesystem contents.

System	HFS Support	HFS+ Support
Linux	Kernel module `hfs`	Kernel module `hfsplus`
FreeBSD	Native support	Native support
macOS	Native support	Native support
OpenBSD	Via FUSE	Via FUSE
NetBSD	Via FUSE	Via FUSE

Table 2.1: Platform Support Matrix

2.1.3 Key Features

Full Specification Compliance

All utilities strictly adhere to:

- Inside Macintosh: Files (HFS specification)

- Apple Technical Note TN1150 (HFS+ specification)

- Unix/POSIX filesystem utility standards

Zero-Tolerance Validation

The test suite implements a "zero-tolerance" policy:

- ANY deviation from specification = test failure

- ALL filesystems must be 100% correct

- Complete validation before and after fsck operations

Journaling Support

HFS+ journaling is supported with appropriate warnings:

- Journal creation with `mkfs.hfs+ -j`

- Journal validation and replay in `fsck.hfs+`

- Linux kernel compatibility warnings (journaling not supported in Linux HFS+ driver)

Date Limit Awareness

The utilities handle filesystem date limits correctly:

- HFS: Maximum date February 6, 2028 (Y2K28 problem)

- HFS+: Maximum date February 6, 2040 (Y2K40 problem)

- Automatic date correction to safe values

2.2 Installation

2.2.1 Building from Source

Prerequisites

```
1  # Debian/Ubuntu
2  sudo apt-get install build-essential
3
4  # Fedora/RHEL
5  sudo dnf install gcc make
6
7  # macOS
8  xcode-select --install
9
10 # BSD
11 # Compiler included by default
```

Compilation

```
1  # Clone repository
2  git clone https://github.com/JotaRandom/hfsutils.git
3  cd hfsutils
4
5  # Build all utilities
6  make
7
8  # Build specific sets
9  make set-hfs        # mkfs.hfs, fsck.hfs, mount.hfs
10 make set-hfsplus    # mkfs.hfs+, fsck.hfs+, mount.hfs+
11
12 # Build with hfsutil
13 make all
```

Installation Options

Linux systems (with kernel HFS/HFS+ drivers):

```
1  sudo make install-linux PREFIX=/usr
```

Complete installation (filesystem utilities + hfsutil):

```
1  sudo make install-complete PREFIX=/usr/local
```

Individual utilities:

```
1  sudo make install-mkfs.hfs+ PREFIX=/usr
2  sudo make install-fsck.hfs PREFIX=/usr
3  sudo make install-mount.hfs+ PREFIX=/usr
```

2.2.2 Verifying Installation

After installation, verify the utilities are available:

```
1  mkfs.hfs --version
2  mkfs.hfs+ --version
3  fsck.hfs --version
```

```
4  fsck.hfs+ --version
5  mount.hfs --help
6  mount.hfs+ --help
7  hfsutil --version    # If installed
```

Check manpages:

```
1  man mkfs.hfs
2  man mkfs.hfs+
3  man fsck.hfs+
4  man mount.hfs+
```

2.3 Quick Start

2.3.1 Creating a Filesystem

Create an HFS filesystem:

```
1  # Create 10MB image file
2  dd if=/dev/zero of=test.img bs=1M count=10
3
4  # Format as HFS
5  mkfs.hfs -L "MyDisk" test.img
```

Create an HFS+ filesystem:

```
1  # Create 50MB image file
2  dd if=/dev/zero of=test.img bs=1M count=50
3
4  # Format as HFS+
5  mkfs.hfs+ -L "MyDisk" test.img
6
7  # Format with journaling (Linux warning will appear)
8  mkfs.hfs+ -j -L "MyDisk" test.img
```

2.3.2 Checking a Filesystem

Verify filesystem integrity:

```
1  # Check HFS filesystem (read-only)
2  fsck.hfs -n test.img
3
4  # Check and repair HFS+ filesystem
5  fsck.hfs+ -y test.img
6
7  # Verbose output
8  fsck.hfs+ -v -n test.img
```

2.3.3 Mounting a Filesystem

On systems with kernel HFS/HFS+ support:

```
1  # Create mount point
2  sudo mkdir /mnt/hfs
3
4  # Mount read-write
5  sudo mount.hfs+ test.img /mnt/hfs
```

```
6
7   # Mount read-only
8   sudo mount.hfs+ -r test.img /mnt/hfs
9
10  # Unmount
11  sudo umount /mnt/hfs
```

2.3.4   Using hfsutil Commands

On systems without kernel support or for direct access:

```
1   # Format volume
2   hfsutil hformat -L "MyDisk" test.img
3
4   # Mount in hfsutil
5   hfsutil hmount test.img
6
7   # List contents
8   hfsutil hls
9
10  # Copy file into volume
11  hfsutil hcopy myfile.txt :myfile.txt
12
13  # Copy file out of volume
14  hfsutil hcopy :myfile.txt retrieved.txt
15
16  # Unmount
17  hfsutil humount
```

2.4   Documentation Structure

This manual is organized as follows:

Chapter 2: HFS Specification Details of the classic HFS filesystem format

Chapter 3: HFS+ Specification Details of the HFS+ filesystem format with journaling

Chapter 4: mkfs Utilities Filesystem creation tools

Chapter 5: fsck Utilities Filesystem checking and repair tools

Chapter 6: mount Utilities Filesystem mounting tools

Chapter 7: hfsutil Commands File manipulation without mounting

Chapter 8: Implementation Details Internal architecture and algorithms

Chapter 9: Testing and Validation Test suite and quality assurance

Chapter 10: Appendix Structure definitions, error codes, glossary

Chapter 3

HFS Specification

3.1 HFS Filesystem Overview

The Hierarchical File System (HFS) is the filesystem used by Apple Computer for Mac OS systems from 1985 until Mac OS X. Also known as "Mac OS Standard" or "HFS Classic", it provides a hierarchical directory structure with support for file and folder metadata.

3.1.1 Key Characteristics

- **Maximum volume size**: 2 TB (2,199,023,255,552 bytes)

- **Maximum file size**: 2 GB (2,147,483,647 bytes)

- **Filename length**: 31 bytes (Macintosh Roman encoding)

- **Date range**: January 1, 1904 to February 6, 2028 (Y2K28 limit)

- **Allocation block size**: 512 bytes minimum, 64 KB maximum

- **Maximum allocation blocks**: 65,535 (16-bit addressing)

- **Case sensitivity**: Case-insensitive, case-preserving

- **Byte order**: Big-endian (MSB first)

3.1.2 Volume Structure

An HFS volume is divided into **logical blocks** (512 bytes each) and **allocation blocks** (multiples of logical blocks).

Complete Volume Layout

Offset (bytes)	Size	Description
0	1024	Boot blocks (2 logical blocks)
1024	512	Master Directory Block (MDB)
1536	Variable	Allocation bitmap start
...	Variable	Extents B-tree file
...	Variable	Catalog B-tree file
...	Variable	File data area
volume_size - 1536	512	(Reserved space)
volume_size - 1024	512	Alternate MDB

Offset (bytes)	Size	Description
volume_size - 512	512	(Last logical block)

Table 3.1: HFS Physical Volume Layout

Critical Formula for Alternate MDB:

$$\text{Alt\_MDB\_offset} = \text{total\_bytes} - 1024 \tag{3.1}$$

This is **NOT** (`total_sectors - 2) * 512`. It is literally 1024 bytes before the end, regardless of sector size.

3.2 Master Directory Block (MDB) - Complete Specification

The MDB is the **single most critical structure** in HFS. It is located at byte offset 1024 and is exactly 162 bytes long.

3.2.1 MDB Complete Field Map - Every Byte Documented

Offset	Field Name	Type	Bytes	Description
+0	drSigWord	uint16	2	Signature: **0x4244** ('BD')
+2	drCrDate	uint32	4	Creation date (Mac time)
+6	drLsMod	uint32	4	Last modification date
+10	drAtrb	uint16	2	Volume attributes (see below)
+12	drNmFls	uint16	2	Files in root directory
+14	drVBMSt	uint16	2	First allocation bitmap block
+16	drAllocPtr	uint16	2	Start of next allocation search
+18	drNmAlBlks	uint16	2	Number of allocation blocks
+20	drAlBlkSiz	uint32	4	Allocation block size (bytes)
+24	drClpSiz	uint32	4	Default clump size
+28	drAlBlSt	uint16	2	First allocation block
+30	drNxtCNID	uint32	4	Next Catalog Node ID
+34	drFreeBks	uint16	2	Free allocation blocks
+36	drVN	Pstring	28	Volume name (1 len + 27 chars)
+64	drVolBkUp	uint32	4	Last backup date
+68	drVSeqNum	uint16	2	Backup sequence number
+70	drWrCnt	uint32	4	Volume write count
+74	drXTClpSiz	uint32	4	Extents clump size
+78	drCTClpSiz	uint32	4	Catalog clump size
+82	drNmRtDirs	uint16	2	Directories in root
+84	drFilCnt	uint32	4	Total files on volume
+88	drDirCnt	uint32	4	Total directories
+92	drFndrInfo	uint32[8]	32	Finder information
+124	drVCSize	uint16	2	Volume cache size
+126	drVBMCSize	uint16	2	Bitmap cache size
+128	drCtlCSize	uint16	2	Common cache size
+130	drXTFlSize	uint32	4	Extents file size
+134	drXTExtRec	ExtRec[3]	12	Extents file extents

Offset	Field Name	Type	Bytes	Description
+146	drCTFlSize	uint32	4	Catalog file size
+150	drCTExtRec	ExtRec[3]	12	Catalog file extents

Table 3.2: Master Directory Block Complete Structure (162 bytes total)

3.2.2 Critical Field Details - Bit-by-Bit

drSigWord - Signature (Offset +0, 2 bytes)

Value: 0x4244 (big-endian)
   Byte representation:

```
Offset 1024: 0x42 ('B')
Offset 1025: 0x44 ('D')
```

   Validation:

- Must be exactly 0x4244

- Any other value = not HFS or corrupted

- Endianness test: if you read 0x4442, you're reading little-endian

   Oddity: The origin of "BD" is unknown. Speculation includes "Block Device" but Apple never documented it.

drAtrb - Volume Attributes (Offset +10, 2 bytes)

16-bit flags field (big-endian). **Every bit documented**:

Bit	Hex Mask	Meaning
0-6	0x007F	Reserved (must be 0)
7	0x0080	Volume locked by hardware
8	**0x0100**	**Volume unmounted properly**
9	0x0200	Volume has spared bad blocks
10	0x0400	Volume needs consistency check (kNeedRebuild)
11	0x0800	Catalog node IDs reused (kBadCNID)
12	0x1000	Unused nodes fix needed
13	0x2000	Volume journaled (HFS+ only, not used in HFS)
14	0x4000	Software lock
15	0x8000	Spare boot blocks (not used)

Table 3.3: drAtrb Bit Definitions

   Critical: Bit 8 (0x0100) MUST be set for clean unmount. mkfs.hfs sets:

```
drAtrb = 0x0100  // Big-endian bytes: 0x01 0x00
```

   Hex dump verification:

```
xxd -s 1034 -l 2 -p volume.hfs
Expected output: 0100
```

drNxtCNID - Next Catalog Node ID (Offset +30, 4 bytes)

Value: 32-bit unsigned integer, big-endian
    Minimum: 0x00000010 (16 decimal)
    Reserved CNIDs 1-15:

CNID	Purpose
1	Parent of root directory (kHFSRootParentID)
2	Root directory (kHFSRootFolderID)
3	Extents overflow file (kHFSExtentsFileID)
4	Catalog file (kHFSCatalogFileID)
5	Bad allocation blocks file (kHFSBadBlock-FileID)
6-15	Reserved, not used in HFS

Table 3.4: Reserved Catalog Node IDs

    Byte representation for value 16:

```
Offset 1054: 0x00
Offset 1055: 0x00
Offset 1056: 0x00
Offset 1057: 0x10
```

    Common error: If you see 0x00000000, the MDB was not initialized correctly. The volume cannot be used.

drVN - Volume Name (Offset +36, 28 bytes)

Format: Pascal string (length-prefixed)

- Byte 0: Length (0-27)

- Bytes 1-27: Characters (Macintosh Roman encoding)

    Example: "MyDisk"

```
Offset 1060: 0x06           // Length = 6
Offset 1061: 0x4D ('M')
Offset 1062: 0x79 ('y')
Offset 1063: 0x44 ('D')
Offset 1064: 0x69 ('i')
Offset 1065: 0x73 ('s')
Offset 1066: 0x6B ('k')
Offset 1067-1087: 0x00     // Padding
```

    Restrictions:

- Length: 1-27 bytes (0 = invalid)

- Cannot contain colon (:) - path separator

- Macintosh Roman encoding (not UTF-8!)

    Oddity: Unlike HFS+, HFS uses Pascal strings (length prefix) instead of C strings (null-terminated).

3.2.3 Extent Records - Complete Structure

An extent record describes up to 3 contiguous runs of allocation blocks.

Extent Descriptor (4 bytes each)

Offset	Field	Description
+0	startBlock	First allocation block (uint16)
+2	blockCount	Number of blocks (uint16)

Table 3.5: Extent Descriptor Structure

Extent Record (12 bytes)

3 consecutive extent descriptors:

```
ExtentRecord {
    ExtentDescriptor[0]:  // Bytes 0-3
        uint16 startBlock
        uint16 blockCount
    ExtentDescriptor[1]:  // Bytes 4-7
        uint16 startBlock
        uint16 blockCount
    ExtentDescriptor[2]:  // Bytes 8-11
        uint16 startBlock
        uint16 blockCount
}
```

Unused extents: Set both fields to 0.
Example: File uses blocks 100-109 and 200-249:

```
Extent[0]: startBlock=100, blockCount=10
Extent[1]: startBlock=200, blockCount=50
Extent[2]: startBlock=0,   blockCount=0   // Unused
```

3.2.4 Alternate MDB - Critical Backup

The alternate MDB is an **exact copy** of the primary MDB.

Location Calculation

Precise formula:
$$\text{alt\_offset} = \text{device\_size\_bytes} - 1024 \tag{3.2}$$

Example for 10 MB volume:

```
Device size: 10,485,760 bytes
Alt MDB at:  10,485,760 - 1,024 = 10,484,736 bytes
```

Verification:

```
FILESIZE=$(stat -c%s volume.hfs)
ALTOFFSET=$((FILESIZE - 1024))
xxd -s $ALTOFFSET -l 2 -p volume.hfs
# Expected: 4244 (same signature as primary)
```

Common mistake: Using (`num_sectors - 2) * 512` assumes 512-byte sectors. The specification is **always** "1024 bytes before end", regardless of sector size.

3.3   HFS B-Trees - Complete Specification

3.3.1   B-Tree Overview

HFS uses B-trees for the catalog (files/folders) and extents overflow (fragmented files).

Key Characteristics

- **Node size**: Fixed 512 bytes for HFS

- **Balancing**: Self-balancing tree structure

- **Key comparison**: Case-insensitive for catalog names

- **Depth**: Typically 2-4 levels for most volumes

3.3.2   Node Descriptor - First 14 Bytes of Every Node

Every B-tree node starts with this 14-byte header:

Offset	Field	Description
+0	fLink	Forward link: next node at this level (uint32)
+4	bLink	Backward link: previous node (uint32)
+8	kind	Node type (int8, see below)
+9	height	Node level: 0 = leaf, 1+ = index (uint8)
+10	numRecords	Number of records in node (uint16)
+12	reserved	Reserved, must be 0 (uint16)

Table 3.6: Node Descriptor (14 bytes)

Node Types (kind field)

Value	Meaning
-1 (0xFF)	Index node (internal)
0	Header node (node 0only)
1	Map node (allocation bitmap)
2	Leaf node (data records)

Table 3.7: B-Tree Node Types

Oddity: Index nodes use -1 (signed), not 255 (unsigned). This is a signed int8 field.

3.3.3   Header Node (Node 0) - Complete Structure

The first node (offset 0 in B-tree file) is always the header node.

BTHeaderRec Structure (106 bytes)

Located at offset 14 (after node descriptor):

Offset	Field	Type	Description
+14	treeDepth	uint16	Current depth (0 = empty)
+16	rootNode	uint32	Node number of root
+20	leafRecords	uint32	Total leaf records
+24	firstLeafNode	uint32	First leaf node number
+28	lastLeafNode	uint32	Last leaf node number
+32	nodeSize	uint16	Node size (**512 for HFS**)
+34	maxKeyLength	uint16	Maximum key length
+36	totalNodes	uint32	Total nodes in tree
+40	freeNodes	uint32	Unused nodes
+44	reserved1	uint16	Reserved
+46	clumpSize	uint32	Clump size (bytes)
+50	btreeType	uint8	0=Catalog, 255=Extents
+51	reserved2	uint8	Reserved
+52	attributes	uint32	B-tree attributes
+56	reserved3	uint8[64]	Reserved

Table 3.8: BTHeaderRec Structure

Critical: nodeSize MUST be 512 for HFS. HFS+ uses 4096, but HFS is fixed at 512.

3.3.4 Date/Time Representation

HFS uses Mac absolute time: unsigned 32-bit seconds since midnight, January 1, 1904 GMT.
Conversion formulas:

```
MAC_EPOCH = 2082844800  // Offset from Unix epoch

Unix to Mac: mac_time = unix_time + MAC_EPOCH
Mac to Unix: unix_time = mac_time - MAC_EPOCH
```

Y2K28 Problem:

$$\text{Max\_date} = 1904 + \frac{2^{32}}{365.25 \times 24 \times 3600} \approx 2028 \tag{3.3}$$

Specifically: February 6, 2028, 06:28:15 GMT
Safe date handling: hfsutils uses `hfs_get_safe_time()` which caps dates at 2028.

3.4 Byte Order (Endianness) - Critical

All HFS multi-byte fields are big-endian.

3.4.1 Endianness Examples

16-bit value 0x1234:

```
Byte 0: 0x12 (MSB)
Byte 1: 0x34 (LSB)
```

32-bit value 0x12345678:

```
Byte 0: 0x12 (Most significant)
Byte 1: 0x34
Byte 2: 0x56
Byte 3: 0x78 (Least significant)
```

Writing code:

```
// WRONG - host byte order
uint16_t value = 0x4244;
write(fd, &value, 2);  // Writes 0x44 0x42 on little-endian!

// CORRECT - explicit byte order
unsigned char sig[2] = {0x42, 0x44};
write(fd, sig, 2);     // Always correct
```

3.5 Oddities, Edge Cases, and Implementation Notes

3.5.1 The 16-Bit Limitation

HFS uses 16-bit allocation block numbers, limiting volumes to 65,535 blocks maximum.
Volume size calculation:

$$\text{max\_volume} = 65535 \times \text{block\_size} \tag{3.4}$$

For 32 KB blocks:

$$65535 \times 32768 = 2,147,450,880 \text{ bytes} \approx 2 \text{ GB} \tag{3.5}$$

Limitation: Cannot create HFS volumes larger than ˜2 TB (with 64 KB blocks).

3.5.2 Pascal Strings vs C Strings

Pascal string (HFS): Length byte + data

```
"\x06Hello!"  // Byte 0 = length, followed by data
```

C string: Null-terminated

```
"Hello!\0"    // Ends with null byte
```

Gotcha: A Pascal string of length 0 is valid (empty string). A C string must have at least the null terminator.

3.5.3 Allocation Block Alignment

File data MUST start on allocation block boundaries. You cannot start a file in the middle of a block.
Implication: Small files waste space. If block size is 4 KB, a 1-byte file wastes 4095 bytes.

3.5.4   MacRoman Character Encoding

HFS uses MacRoman, not ASCII or UTF-8.
    Differences from ASCII:

- Characters 0-127: Same as ASCII

- Characters 128-255: Different from ISO-8859-1

Example oddities:

- 0xD0 = en dash (—) in MacRoman, in ISO-8859-1

- 0xD1 = em dash (—) in MacRoman, $\tilde{N}$ in ISO-8859-1

Implementation: For maximum compatibility, restrict filenames to ASCII 32-126.

Chapter 4

HFS+ Specification

4.1 HFS+ Filesystem Overview

HFS Plus (HFS+), also known as Mac OS Extended, is Apple's modern filesystem designed to replace HFS. Introduced in Mac OS 8.1 (1998), it addresses HFS limitations while maintaining backward compatibility.

4.1.1 Key Improvements over HFS

- **Unicode filenames**: Full Unicode support (up to 255 UTF-16 characters)

- **Larger volumes**: Up to 8 EB (exabytes) theoretical

- **Larger files**: Up to 8 EB per file

- **Smaller allocation blocks**: More efficient space usage

- **Journaling support**: Optional transaction journal for crash recovery

- **Hard links**: Support for hard links (Mac OS X 10.2+)

- **Symbolic links**: Support for symbolic links

- **Extended attributes**: Arbitrary metadata on files/folders

- **Case sensitivity option**: HFSX variant supports case-sensitive names

4.1.2 HFS+ Characteristics

Feature	Specification
Maximum volume size	8 EB (2^{63} bytes)
Maximum file size	8 EB (2^{63} bytes)
Filename length	255 Unicode characters (UTF-16)
Date range	January 1, 1904 to February 6, 2040 (Y2K40)
Minimum allocation block	512 bytes
Block addressing	32-bit allocation block numbers
Case sensitivity	Case-insensitive (HFS+) or case-sensitive (HFSX)

Table 4.1: HFS+ Specifications

4.2 Volume Structure

HFS+ volumes are structured similarly to HFS but with enhanced metadata structures.

4.2.1 Volume Layout

Location	Name	Description
Byte 0	Reserved	Boot sector (512 bytes)
Byte 512	Reserved	Additional boot area (512 bytes)
Byte 1024	Volume Header	Primary volume metadata
...	Allocation Bitmap	Volume space allocation map
...	Allocation File	B-tree of allocation bitmap extents
...	Extents Overflow File	B-tree of file extent records
...	Catalog File	B-tree of all files and folders
...	Attributes File	B-tree of extended attributes
...	Startup File	Boot loader for non-Mac systems
...	Data Area	File contents and special files
-1024	Alternate VH	Backup copy of Volume Header

Table 4.2: HFS+ Volume Layout

4.3 Volume Header - Complete Bit-Level Specification

The Volume Header is the **most critical structure** in HFS+. Located at byte offset 1024, it is exactly **512 bytes**.

4.3.1 Volume Header Complete Field Map - Every Byte Documented

Offset	Field	Type	Size	Description
+0	signature	uint16	2	**0x482B** ('H+') or **0x4858** ('HX')
+2	version	uint16	2	**4** (HFS+) or **5** (HFSX)
+4	attributes	uint32	4	Volume attributes (flags, see below)
+8	lastMountedVersion	uint32	4	OS signature that last mounted
+12	journalInfoBlock	uint32	4	Journal info block (0 = no journal)
+16	createDate	uint32	4	Creation date (HFS+ time)
+20	modifyDate	uint32	4	Last modification date
+24	backupDate	uint32	4	Last backup date
+28	checkedDate	uint32	4	Last fsck date
+32	fileCount	uint32	4	Total files on volume
+36	folderCount	uint32	4	Total folders on volume
+40	blockSize	uint32	4	Allocation block size (bytes)
+44	totalBlocks	uint32	4	Total allocation blocks
+48	freeBlocks	uint32	4	Free allocation blocks
+52	nextAllocation	uint32	4	Hint for next allocation

Offset	Field	Type	Size	Description
+56	rsrcClumpSize	uint32	4	Default resource fork clump
+60	dataClumpSize	uint32	4	Default data fork clump
+64	nextCatalogID	uint32	4	Next unused Catalog Node ID
+68	writeCount	uint32	4	Volume write count
+72	encodingsBitmap	uint64	8	Text encodings used (64 bits)
+80	finderInfo	uint32[8]	32	Finder information
+112	allocationFile	HFSPlusForkData	80	Allocation file fork
+192	extentsFile	HFSPlusForkData	80	Extents file fork
+272	catalogFile	HFSPlusForkData	80	Catalog file fork
+352	attributesFile	HFSPlusForkData	80	Attributes file fork
+432	startupFile	HFSPlusForkData	80	Startup file fork

Table 4.3: HFS+ Volume Header Structure (512 bytes total)

Total size verification: $2 + 2 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 8 + 32 + (80*5) = 512$ bytes

4.3.2 Critical Field Details - Bit-by-Bit

signature - Volume Signature (Offset +0, 2 bytes)

HFS+ Standard: 0x482B (big-endian)
  Byte representation:

```
Offset 1024: 0x48 ('H')
Offset 1025: 0x2B ('+')
```

  HFSX Case-Sensitive: 0x4858 (big-endian)
  Byte representation:

```
Offset 1024: 0x48 ('H')
Offset 1025: 0x58 ('X')
```

  Validation:

- Must be exactly 0x482B or 0x4858

- Any other value = not HFS+ or corrupted

- 0x4244 = HFS (not HFS+)

  Hex dump verification:

```
xxd -s 1024 -l 2 -p volume.hfsplus
Expected: 482b (HFS+) or 4858 (HFSX)
```

version - Volume Version (Offset +2, 2 bytes)

HFS+ Standard: 0x0004 (4 decimal, big-endian)
  Byte representation:

```
Offset 1026: 0x00
Offset 1027: 0x04
```

  HFSX: 0x0005 (5 decimal)
  Validation: mkfs.hfs+ sets version to 4 for standard HFS+.

attributes - Volume Attributes (Offset +4, 4 bytes)

32-bit flags field (big-endian). **Every bit documented**:

Bit	Hex Mask	Meaning
0-6	0x0000007F	Reserved (must be 0)
7	0x00000080	Volume locked by hardware
8	**0x00000100**	**Volume unmounted cleanly**
9	0x00000200	Volume has spared bad blocks
10	0x00000400	Needs fsck (consistency check required)
11	0x00000800	Catalog node IDs wrapped around
12	0x00001000	Software lock
13	**0x00002000**	**Volume is journaled**
14	0x00004000	Reserved
15	0x00008000	Reserved
16-31	0xFFFF0000	Reserved

Table 4.4: Volume Header attributes Bit Definitions

Common values:

- 0x00000100: Clean, non-journaled (mkfs.hfs+ default)

- 0x00002100: Clean, journaled (mkfs.hfs+ -j)

Byte representation for 0x00000100:

```
Offset 1028: 0x00
Offset 1029: 0x00
Offset 1030: 0x01
Offset 1031: 0x00
```

Verification:

```
xxd -s 1028 -l 4 -p volume.hfsplus
Expected: 00000100 (non-journaled) or 00002100 (journaled)
```

blockSize - Allocation Block Size (Offset +40, 4 bytes)

Value: 32-bit unsigned, big-endian
Valid range:

- Minimum: 512 bytes

- Maximum: Typically 32 KB, theoretically larger

- Must be power of 2

- Must be multiple of 512

Example for 4096 bytes (4 KB):

```
Decimal: 4096
Hex: 0x00001000
Bytes at offset 1064:
  0x00 0x00 0x10 0x00
```

Validation:

```
xxd -s 1064 -l 4 -p volume.hfsplus
# For 4 KB blocks: 00001000
# For 8 KB blocks: 00002000
```

rsrcClumpSize and dataClumpSize (Offsets +56, +60)

Critical: These MUST be non-zero in valid HFS+ volumes.
   Recommended value:
$$\text{clumpSize} = \text{blockSize} \times 4 \tag{4.1}$$

   For 4 KB blocks:

$$\text{clumpSize} = 4096 \times 4 = 16384 \text{ bytes} = \texttt{0x00004000} \tag{4.2}$$

   Byte representation:

```
rsrcClumpSize at offset 1080: 0x00 0x00 0x40 0x00
dataClumpSize at offset 1084: 0x00 0x00 0x40 0x00
```

   Common error: If these are 0x00000000, the Volume Header is invalid and nextCatalogID appears at wrong offset.
   Verification:

```
xxd -s 1080 -l 4 -p volume.hfsplus  # rsrcClumpSize
xxd -s 1084 -l 4 -p volume.hfsplus  # dataClumpSize
# Both should be non-zero
```

nextCatalogID - Next CNID (Offset +64, 4 bytes)

Minimum: 0x00000010 (16 decimal)
   Reserved CNIDs 1-15:

CNID	Purpose
1	Root folder's parent (kHFSRootParentID)
2	Root folder (kHFSRootFolderID)
3	Extents overflow file (kHFSExtentsFileID)
4	Catalog file (kHFSCatalogFileID)
5	Bad allocation blocks file (kHFSBadBlock-FileID)
6	Allocation bitmap file (kHFSAllocationFileID)
7	Startup file (kHFSStartupFileID)
8	Attributes file (kHFSAttributesFileID)
9-13	Reserved
14	Journal file (kHFSJournalFileID, if journaled)
15	Journal info block (kHFSJournalInfoBlockID)

Table 4.5: Reserved HFS+ Catalog Node IDs

   Byte representation for value 16:

```
Offset 1088: 0x00
Offset 1089: 0x00
Offset 1090: 0x00
Offset 1091: 0x10
```

Critical validation:

```
xxd -s 1088 -l 4 -p volume.hfsplus
Expected: 00000010 (minimum value)
```

Common error: If you see 00000000, rsrcClumpSize/dataClumpSize were likely omitted, shifting all subsequent fields.

4.3.3 HFSPlusForkData Structure - 80 Bytes Per Fork

Each fork descriptor is 80 bytes:

Offset	Field	Description
+0	logicalSize	File size in bytes (uint64)
+8	clumpSize	Clump size for this file (uint32)
+12	totalBlocks	Total allocation blocks (uint32)
+16	extents[0]	First extent descriptor (8 bytes)
+24	extents[1]	Second extent descriptor (8 bytes)
+32	extents[2]	Third extent descriptor (8 bytes)
+40	extents[3]	Fourth extent descriptor (8 bytes)
+48	extents[4]	Fifth extent descriptor (8 bytes)
+56	extents[5]	Sixth extent descriptor (8 bytes)
+64	extents[6]	Seventh extent descriptor (8 bytes)
+72	extents[7]	Eighth extent descriptor (8 bytes)

Table 4.6: HFSPlusForkData Structure (80 bytes)

HFSPlusExtentDescriptor - 8 Bytes Each

Offset	Field	Description
+0	startBlock	First allocation block (uint32)
+4	blockCount	Number of blocks (uint32)

Table 4.7: HFSPlusExtentDescriptor (8 bytes)

Example: Catalog file uses blocks 100-199:

```
logicalSize:  0x0000000000018800   (100 KB)
clumpSize:    0x00004000           (16 KB)
totalBlocks:  0x00000019           (25 blocks)
extents[0]:   startBlock=100, blockCount=25
              Bytes: 00 00 00 64 00 00 00 19
extents[1-7]: startBlock=0, blockCount=0 (unused)
```

4.3.4 Alternate Volume Header Location

Exact formula:
$$\text{alt\_VH\_offset} = \text{volume\_size\_bytes} - 1024 \qquad (4.3)$$

Same as HFS. This is **NOT** sector-based.
Example for 50 MB:

```
Volume size: 52,428,800 bytes
Alt VH at:   52,428,800 - 1,024 = 52,427,776 bytes
```

Verification:

```
FILESIZE=$(stat -c%s volume.hfsplus)
ALTOFFSET=$((FILESIZE - 1024))
xxd -s $ALTOFFSET -l 2 -p volume.hfsplus
Expected: 482b (same as primary)
```

4.4   HFS+ B-Trees - Complete Node and Record Structures

HFS+ uses B-trees for **all metadata organization**: Catalog, Extents, Attributes.

4.4.1   B-Tree Node Structure - 14-Byte Node Descriptor

Every B-tree node begins with a 14-byte descriptor:

Offset	Field	Type	Size	Description
+0	fLink	uint32	4	Forward link (next node, 0 = none)
+4	bLink	uint32	4	Backward link (prev node, 0 = none)
+8	kind	int8	1	Node type (see below)
+9	height	uint8	1	Level in tree (0 = leaf)
+10	numRecords	uint16	2	Number of records in node
+12	reserved	uint16	2	Reserved (must be 0)

Table 4.8: BTNodeDescriptor (14 bytes)

kind - Node Type Values

Value	Node Type
-1	Leaf node (contains data records)
0	Index node (contains child pointers)
1	Header node (B-tree metadata, always node 0)
2	Map node (allocation bitmap)

Table 4.9: B-tree Node Types

Byte example for header node (kind=1):

```
Offset +8: 0x01 (header node)
Offset +9: 0x00 (height 0, not used in header)
Offset +10-11: 0x00 0x03 (3 records typical: header, user data, map)
```

4.4.2   BTHeaderRec - B-Tree Header Record (106 Bytes)

Located in the **first record of node 0 (header node)** in every B-tree.

Offset	Field	Type	Size	Description
+0	treeDepth	uint16	2	Current depth (1 = only root)
+2	rootNode	uint32	4	Root node number
+6	leafRecords	uint32	4	Total leaf records
+10	firstLeafNode	uint32	4	First leaf node number
+14	lastLeafNode	uint32	4	Last leaf node number
+18	nodeSize	uint16	2	Node size in bytes
+20	maxKeyLength	uint16	2	Maximum key length
+22	totalNodes	uint32	4	Total nodes allocated
+26	freeNodes	uint32	4	Number of free nodes
+30	reserved1	uint16	2	Reserved
+32	clumpSize	uint32	4	Clump size for B-tree file
+36	btreeType	uint8	1	B-tree type (0=HFS, 128=HFS+)
+37	keyCompareType	uint8	1	Key comparison type
+38	attributes	uint32	4	B-tree attributes (flags)
+42	reserved3	uint32[16]	64	Reserved (must be 0)

Table 4.10: BTHeaderRec Structure (106 bytes total)

Total size verification: $2 + 4 + 4 + 4 + 4 + 2 + 2 + 4 + 4 + 2 + 4 + 1 + 1 + 4 + 64$ = 106 bytes

Critical BTHeaderRec Field Details

treeDepth (Offset +0):

- Value 0: Empty tree

- Value 1: Only root node (contains data directly)

- Value 2+: Root is index node

- **New volume**: Usually 1 (minimal tree)

nodeSize (Offset +18):

- **HFS+ Standard**: 4096 bytes (4 KB)

- **HFS+ Large**: 8192 bytes (8 KB)

- Must be power of 2

- Must be $\geq$ 512 bytes

Byte representation for 4096:

```
Offset +18: 0x10
Offset +19: 0x00
(Big-endian: 0x1000 = 4096)
```

btreeType (Offset +36):

- **0**: HFS B-tree (legacy)

- **128**: HFS+ B-tree (standard)

- **255**: Reserved

keyCompareType (Offset +37):

- **0xBC**: Case-insensitive (HFS+ default)

- **0xCF**: Binary compare (HFSX case-sensitive)

attributes (Offset +38, 4 bytes):

Bit	Hex Mask	Meaning
0	0x00000001	Bad close (B-tree not closed properly)
1	0x00000002	Big keys (key length > 255 bytes)
2	0x00000004	Variable index keys
3-31	0xFFFFFFF8	Reserved (must be 0)

Table 4.11: BTHeaderRec attributes Flags

4.4.3 Catalog File B-Tree - Complete Key and Record Formats

The Catalog File contains **all files and folders**. It uses **Unicode HFSUniStr255 keys**.

HFSPlusCatalogKey - Variable Length

Offset	Field	Type	Size	Description
+0	keyLength	uint16	2	Total key length (excluding this field)
+2	parentID	uint32	4	Parent folder CNID
+6	nodeName	HFSUniStr255	var	Unicode filename (see below)

Table 4.12: HFSPlusCatalogKey Structure

HFSUniStr255 - Unicode String (Max 255 UTF-16 Characters)

Offset	Field	Description
+0	length	uint16: Number of UTF-16 characters (NOT bytes)
+2	unicode	uint16[]: UTF-16BE characters

Table 4.13: HFSUniStr255 Structure

Example: Filename "Test.txt" (8 characters)

```
length:  0x0008 (8 UTF-16 chars)
unicode:
  0x0054 ('T')
  0x0065 ('e')
  0x0073 ('s')
  0x0074 ('t')
  0x002E ('.')
```

```
  0x0074 ('t')
  0x0078 ('x')
  0x0074 ('t')
Total: 2 + (8*2) = 18 bytes for nodeName
```

Complete catalog key for "Test.txt" in root (CNID 2):

```
keyLength: 0x0016 (22 bytes: 4 parentID + 18 nodeName)
parentID:  0x00000002 (root folder)
nodeName:  [18 bytes as shown above]
Total key: 2 + 22 = 24 bytes
```

Catalog Record Types - 4 Types

Type	Value	Description
kHFSPlusFolderRecord	0x0001	Folder (directory)
kHFSPlusFileRecord	0x0002	File
kHFSPlusFolderThreadRecord	0x0003	Folder thread (CNID → name)
kHFSPlusFileThreadRecord	0x0004	File thread (CNID → name)

Table 4.14: Catalog Record Type Values

HFSPlusCatalogFile - File Record (248 Bytes)

Offset	Field	Type	Size	Description
+0	recordType	int16	2	**0x0002** (file)
+2	flags	uint16	2	File flags
+4	reserved1	uint32	4	Reserved
+8	fileID	uint32	4	Catalog Node ID (CNID)
+12	createDate	uint32	4	Creation date (HFS+ time)
+16	contentModDate	uint32	4	Content modification date
+20	attributeModDate	uint32	4	Attribute modification date
+24	accessDate	uint32	4	Last access date
+28	backupDate	uint32	4	Backup date
+32	permissions	HFSPlusBSDInfo	16	BSD ownership/permissions
+48	userInfo	FInfo	16	Finder user info
+64	finderInfo	FXInfo	16	Finder extended info
+80	textEncoding	uint32	4	Text encoding hint
+84	reserved2	uint32	4	Reserved
+88	dataFork	HFSPlusForkData	80	Data fork descriptor
+168	resourceFork	HFSPlusForkData	80	Resource fork descriptor

Table 4.15: HFSPlusCatalogFile Structure (248 bytes total)

Total size verification: 2 + 2 + 4 + 4 + 4 + 4 + 4 + 4 + 4 + 16 + 16 + 16 + 4 + 4 + 80 + 80 = 248 bytes

4.4.4   Catalog File

The Catalog File is a B-tree containing all files and folders on the volume.

Catalog Keys

- **keyLength**: 2 bytes

- **parentID**: 4 bytes (Parent folder CNID)

- **nodeName**: Variable-length Unicode string

Catalog Record Types

1. **Folder Record (0x0001)**: Directory metadata

2. **File Record (0x0002)**: File metadata and fork information

3. **Folder Thread (0x0003)**: Maps CNID to parent folder

4. **File Thread (0x0004)**: Maps file CNID to parent folder

File Record Structure

- recordType: 0x0002

- flags: File flags

- fileID: Catalog Node ID

- createDate, modifyDate: Timestamps

- permissions: Unix permissions

- userInfo, finderInfo: Finder metadata

- textEncoding: Filename encoding hint

- dataFork: Fork data for data fork (80 bytes)

- resourceFork: Fork data for resource fork (80 bytes)

4.4.5   Extents Overflow File

B-tree storing additional extent records when a file's fork exceeds the 8 extents stored in the catalog.

Extent Key

- keyLength: 2 bytes

- forkType: 0x00 (data) or 0xFF (resource)

- fileID: Catalog Node ID

- startBlock: Starting allocation block

Extent Descriptor

- startBlock: Starting allocation block (4 bytes)

- blockCount: Number of contiguous blocks (4 bytes)

4.4.6   Attributes File

B-tree storing extended attributes (metadata) for files and folders.

Supported Attributes

- Extended attributes (xattrs)

- Access Control Lists (ACLs)

- Resource fork data (if not inline)

- Compressed data (HFS+ compression)

1. fsck.hfs+ detects journal

2. Scans journal for uncommitted transactions

3. Replays completed transactions to restore consistency

4. Marks volume clean after successful replay

4.5   Unicode Normalization - CRITICAL for Filename Compatibility

HFS+ uses **Unicode Normalization Form D (NFD)** for all filenames. This is **mandatory** and causes significant compatibility issues with other systems.

4.5.1   NFD vs NFC - The Core Problem

Unicode allows multiple representations of the same character:

- **NFC (Composed)**: Single codepoint for accented characters

- **NFD (Decomposed)**: Base character + combining accent

Example: Letter "é" (e with acute accent)

Form	Representation
NFC	U+00E9 (single codepoint: LATIN SMALL LETTER E WITH ACUTE)
NFD	U+0065 U+0301 (two codepoints: LATIN SMALL LETTER E + COMBINING ACUTE ACCENT)

Table 4.16: Unicode Normalization Example

Byte representation in UTF-16BE:

```
NFC (1 UTF-16 unit):  0x00E9
NFD (2 UTF-16 units): 0x0065 0x0301

In HFSUniStr255:
  length (NFC): 0x0001 (1 character)
  length (NFD): 0x0002 (2 characters)
```

4.5.2 HFS+ NFD Requirement - MANDATORY

Apple Technical Note TN1150: All HFS+ filenames MUST be stored in NFD form.
 Conversion algorithm:

1. Receive filename from user (may be in any form)

2. Decompose to NFD using Unicode decomposition tables

3. Store in catalog with NFD form

4. When reading, return NFD form to user

 Critical implementation detail:

- mkfs.hfs+ must accept filenames and convert to NFD

- Catalog B-tree keys are compared in NFD form

- Case-insensitive comparison uses Unicode case folding tables

4.5.3 Common NFD Characters - Complete Table

Note: The following table shows common accented characters and their NFD decompositions. Character names are given in ASCII descriptions to avoid PDF encoding issues.

Character	NFC	NFD	Components
a with grave	U+00E0	U+0061 U+0300	a + combining grave
a with acute	U+00E1	U+0061 U+0301	a + combining acute
a with circumflex	U+00E2	U+0061 U+0302	a + combining circumflex
a with tilde	U+00E3	U+0061 U+0303	a + combining tilde
a with diaeresis	U+00E4	U+0061 U+0308	a + combining diaeresis
n with tilde	U+00F1	U+006E U+0303	n + combining tilde
c with cedilla	U+00E7	U+0063 U+0327	c + combining cedilla
u with diaeresis	U+00FC	U+0075 U+0308	u + combining diaeresis
o with diaeresis	U+00F6	U+006F U+0308	o + combining diaeresis
a with ring above	U+00E5	U+0061 U+030A	a + combining ring above

Table 4.17: Common NFD Decompositions

4.5.4 Compatibility Issues

Linux/Windows: Use NFC by default
 Problem: Filename created on macOS with accented characters in NFD form appears as different file than same filename in NFC form created on Linux on the same HFS+ volume.
 Example: A filename like "cafe.txt" with e-acute stored as NFD (e + U+0301) appears different than NFC (U+00E9) - these are different catalog entries!
 Workaround: Always normalize to NFD when writing to HFS+.
 Implementation in hfsutils:

```
// Pseudo-code for filename conversion
void normalize_to_nfd(uint16_t *unicode, size_t *length) {
    // For each character:
    //   1. Look up in Unicode decomposition table
    //   2. Replace with base + combining characters
    //   3. Update length accordingly
}
```

4.6 HFS+ Time Format - Mac Epoch and Conversion

HFS+ uses a **32-bit unsigned integer** for all timestamps, representing seconds since the **Mac epoch**.

4.6.1 Mac Epoch Definition

Mac epoch: January 1, 1904 00:00:00 UTC
   Unix epoch: January 1, 1970 00:00:00 UTC
   Difference: 2,082,844,800 seconds (66 years)

4.6.2 Date Range

With 32-bit unsigned integer:

- Minimum: 0 (January 1, 1904)

- Maximum: 4,294,967,295 (February 6, 2040 06:28:15 UTC)

   Y2K40 Problem: HFS+ timestamps overflow on February 6, 2040.

4.6.3 Conversion Formulas

HFS+ to Unix time:
$$\text{unix\_time} = \text{hfs\_time} - 2082844800 \tag{4.4}$$

   Unix to HFS+ time:

$$\text{hfs\_time} = \text{unix\_time} + 2082844800 \tag{4.5}$$

   Example conversion:

```
HFS+ time:  3600000000 (0xD693A400)
Unix time:  3600000000 - 2082844800 = 1517155200
Unix date:  January 28, 2018 16:00:00 UTC
```

4.6.4 Byte Representation

All timestamps are big-endian 32-bit unsigned integers.
   Example: December 25, 2020 12:00:00 UTC

```
Unix timestamp: 1608897600
HFS+ timestamp: 1608897600 + 2082844800 = 3691742400
Hex: 0xDBF49140
Bytes: 0xDB 0xF4 0x91 0x40
```

   Verification in Volume Header createDate (offset +16):

```
xxd -s 1040 -l 4 -p volume.hfsplus
Expected format: DBXXXXXX (for recent dates)
```

4.6.5 Y2K40 Safeguards in hfsutils

Implementation in src/common/hfstime.c:

```
#define HFS_Y2K40_LIMIT 4294967295
#define HFS_SAFE_YEAR_2030 4102444800

uint32_t hfs_get_safe_time(void) {
    time_t now = time(NULL);
    uint32_t hfs_time = (uint32_t)now + 2082844800;

    // If beyond Y2K40, use January 1, 2030
    if (hfs_time > HFS_Y2K40_LIMIT) {
        hfs_time = HFS_SAFE_YEAR_2030;
    }

    return hfs_time;
}
```

Critical: mkfs.hfs, mkfs.hfs+, and fsck.hfs+ all use this function.

4.7 HFS+ Critical Oddities and Edge Cases

4.7.1 Case-Insensitive vs Case-Preserving

HFS+ Standard Behavior:

- **Case-preserving**: Stores "MyFile.txt" as typed

- **Case-insensitive**: "myfile.txt" and "MyFile.txt" are the SAME file

- Uses Unicode case folding for comparison

  HFSX Behavior:

- **Case-sensitive**: "myfile.txt" and "MyFile.txt" are DIFFERENT files

- Signature: 0x4858 ('HX'), version 5

- keyCompareType: 0xCF (binary compare)

  Incompatibility: Standard HFS+ cannot be converted to HFSX without reformatting.

4.7.2 Folder Valence - Hidden Complexity

HFSPlusCatalogFolder structure includes "valence" field:

- Counts number of items in folder

- **Does NOT include invisible files** (e.g., .DS_Store)

- Must be updated on every file creation/deletion

- Inconsistency causes fsck errors

4.7.3 Hard Links - Indirect Nodes

HFS+ supports hard links (multiple names for same file):

- Uses **indirect nodes** with special parent ID

- Hard link parent: 0xFFFFFFFE (reserved)

- Each hard link has unique CNID

- All point to same fileID in hidden directory

Implementation complexity: Requires special catalog traversal logic.

4.7.4 Compression - Undocumented Extension

macOS 10.6+ introduced HFS+ compression (unofficial):

- Compressed data stored in **extended attributes**

- Resource fork contains decompression metadata

- **Not part of original HFS+ spec**

- Third-party implementations typically ignore

4.7.5 Journal Checksum Algorithm - Missing from TN1150

Journal uses CRC32 or similar checksum (not fully documented):

- Checksum in journal header (offset +28)

- Verifies journal integrity before replay

- **Algorithm varies by implementation**

Safe approach: If checksum fails, refuse to replay journal (mount read-only).

4.7.6 Allocation Block Alignment

Critical for performance:

- Allocation blocks should align to physical sectors

- blockSize should be multiple of physical sector size

- Modern drives: 4 KB sectors → use 4 KB allocation blocks

- Misalignment causes read-modify-write penalty

mkfs.hfs+ default: 4096 bytes (optimal for modern drives)

4.7.7 Extended Attributes File - Optional

attributesFile in Volume Header (offset +352):

- Can be empty (logicalSize = 0) on new volumes

- Created on-demand when first extended attribute added

- Uses its own B-tree structure

- Keys: (fileID, attribute name)

Common attributes:

- com.apple.FinderInfo: Finder metadata

- com.apple.ResourceFork: Resource fork data (alternative storage)

- com.apple.decmpfs: Compressed file data

4.8 Reimplementation Checklist - Everything You Need

4.8.1 Data Structures Required

1. Volume Header (512 bytes) - Complete in this document

2. HFSPlusForkData (80 bytes) - Complete in this document

3. HFSPlusExtentDescriptor (8 bytes) - Complete in this document

4. BTNodeDescriptor (14 bytes) - Complete in this document

5. BTHeaderRec (106 bytes) - Complete in this document

6. HFSPlusCatalogKey (variable) - Complete in this document

7. HFSUniStr255 (variable, max 512 bytes) - Complete in this document

8. HFSPlusCatalogFile (248 bytes) - Complete in this document

9. HFSPlusCatalogFolder (88 bytes) - See Apple TN1150

10. JournalInfoBlock (96 bytes) - Complete in this document

4.8.2 Algorithms Required

1. Unicode NFD normalization (use ICU library or tables)

2. Unicode case folding (for case-insensitive comparison)

3. B-tree insertion/deletion (standard CS algorithm)

4. Extent allocation/deallocation

5. Bitmap manipulation (allocation file)

6. CRC32 or checksum (for journal)

7. HFS+ time conversion (formulas in this document)

4.8.3 Validation Commands

All xxd commands in this document can verify:

- Volume signature (offset 1024)

- Volume version (offset 1026)

- Attributes flags (offset 1028)

- blockSize (offset 1064)

- rsrcClumpSize, dataClumpSize (offsets 1080, 1084)

- nextCatalogID (offset 1088)

- Alternate Volume Header (volume_size - 1024)

  fsck.hfs+ validates:

- All B-tree structures

- Folder valence consistency

- Allocation bitmap consistency

- Extent overflow records

- Journal integrity (if present)

4.8.4 No External References Needed

This document contains:

- Every byte offset for critical structures

- All bit flags with hex masks

- Complete formulas for calculations

- Byte examples for verification

- Common error patterns

- Compatibility warnings

  You can reimplement HFS+ with:

1. This chapter (complete specification)

2. Standard Unicode tables (NFD decomposition)

3. Standard B-tree algorithm (CS textbook)

4. CRC32 implementation (standard)

   No internet required after you have these resources.

4.9 HFS+ vs HFSX

HFSX is a variant of HFS+ with case-sensitive filename comparison.

   Note: hfsutils mkfs.hfs+ creates standard HFS+ volumes (case-insensitive). HFSX support is not currently implemented.

Feature	HFS+	HFSX
Signature	0x482B ('H+')	0x4858 ('HX')
Version	4	5
Case sensitivity	No	Yes
Filename comparison	Case-insensitive	Case-sensitive
"File.txt" = "file.txt"	Yes	No

Table 4.18: HFS+ vs HFSX

4.10 Compatibility Considerations

4.10.1 macOS

- Native support for HFS+ and HFSX

- Full journaling support

- All extended attributes supported

- APFS is now default (macOS 10.13+)

4.10.2 Linux

- Kernel module `hfsplus` required

- **No journaling support in kernel driver**

- Basic read/write support

- Some extended attributes supported

- May mount journaled volumes read-only for safety

4.10.3 FreeBSD/OpenBSD/NetBSD

- Native HFS+ read support

- Write support via FUSE

- No journaling support

4.10.4 Windows

- No native HFS+ support

- Third-party drivers available (Paragon, MacDrive)

- Boot Camp drivers provide read-only access

Chapter 5

mkfs Utilities

Chapter 6

Filesystem Creation Utilities

This chapter documents the mkfs.hfs and mkfs.hfs+ utilities for creating HFS and HFS+ filesystems.

6.1   mkfs.hfs - HFS Classic Filesystem Creation

6.1.1   Synopsis

```
mkfs.hfs [-L label] [-l label] [-s size] [-b block_size] device
```

6.1.2   Description

Creates an HFS (Hierarchical File System) Classic volume on the specified device. The utility initializes all required metadata structures including the Master Directory Block (MDB), volume bitmap, extents file, and catalog file.

6.1.3   Options

Option	Description
-L label	Set volume label (primary option, Unix standard)
-l label	Set volume label (alias for -L)
-s size	Explicit volume size in bytes (optional, auto-detected from device)
-b block_size	Allocation block size in bytes (default: auto-calculated, min 512, max 65536)
device	Block device or file to format (required)

Table 6.1: mkfs.hfs Options

6.1.4   Volume Label

Format: Pascal string (1-27 characters)

- First byte: length (1-27)

- Following bytes: MacRoman encoded characters

- Maximum: 27 characters (28 bytes total with length byte)

- Default: "Untitled" if not specified

44

Character restrictions:

- Colon (:) not allowed (path separator)

- MacRoman encoding (not UTF-8)

- ASCII subset recommended for compatibility

6.1.5 Block Size Selection

Automatic algorithm:

```
if (volume_size <= 256 MB)
    block_size = 512 bytes
else if (volume_size <= 512 MB)
    block_size = 1024 bytes
else if (volume_size <= 1 GB)
    block_size = 2048 bytes
else
    block_size = min(32768, optimal_for_size)
```

Constraints:

- Must be power of 2

- Minimum: 512 bytes

- Maximum: 65,536 bytes (64 KB)

- Must be multiple of device sector size

- Total blocks must fit in 16-bit field (max 65,535 blocks)

6.1.6 Initialization Sequence

Step 1: Device Validation

1. Open device for read/write

2. Query device size

3. Verify minimum size (1440 KB for floppy compatibility)

4. Calculate total blocks

5. Verify blocks $\leq$ 65,535 (16-bit limit)

Step 2: Master Directory Block (MDB) Creation

1. Zero 512-byte MDB buffer

2. Set drSigWord = 0x4244 ('BD')

3. Set drCrDate = current HFS time (hfs_get_safe_time())

4. Set drLsMod = current HFS time

5. Set drAtrb = 0x0100 (unmounted cleanly bit)

6. Set drNmFls = 0 (no files initially)

7. Set drVBMSt = 3 (bitmap starts at block 3)

8. Set drAllocPtr = 0 (allocation search starts at 0)

9. Set drNxtCNID = 16 (first available CNID)

10. Set drFreeBks = total_blocks - reserved

11. Set drVN = volume label (Pascal string)

12. Calculate drVBMSt based on allocation block size

13. Initialize extent descriptors for catalog and extents files

Step 3: Allocation Bitmap Initialization

1. Calculate bitmap size: (total_blocks + 7) / 8 bytes

2. Allocate bitmap buffer

3. Mark boot blocks as used (blocks 0-1)

4. Mark MDB as used (block 2)

5. Mark bitmap itself as used

6. Mark catalog file extents as used

7. Mark extents file extents as used

8. Write bitmap to drVBMSt

Step 4: Catalog File Initialization

1. Allocate initial catalog file (default: 10 allocation blocks)

2. Create header node (node 0):

   - Node descriptor: kind = 1 (header)
   - BTHeaderRec: initialize all fields
   - Set bthDepth = 1 (only root node)
   - Set bthRoot = 1 (root is node 1)
   - Set bthNRecs = 0 (no records yet)
   - Set bthFNode = 0, bthLNode = 0

3. Create root folder record (CNID 2) in node 1

4. Write catalog file to allocated extents

Step 5: Extents File Initialization

1. Allocate extents file (default: 5 allocation blocks)

2. Create header node with empty B-tree

3. Initialize BTHeaderRec

4. Set extents file as empty (no overflow extents initially)

Step 6: Alternate MDB

1. Copy primary MDB to buffer

2. Write to offset: device_size_bytes - 1024

3. Verify write succeeded

Step 7: Finalization

1. Sync all writes to device

2. Close device

3. Report success (exit 0)

6.1.7 Exit Codes

Code	Meaning
0	Success, filesystem created
1	Usage error (invalid arguments, missing device) or mkfs failure

Table 6.2: mkfs.hfs Exit Codes (Unix Standard)

6.1.8 Examples

Create HFS volume with label:

```
mkfs.hfs -L "My Volume" /dev/sdb1
mkfs.hfs -L "Backup" disk.hfs
```

Create with specific block size:

```
mkfs.hfs -L "Data" -b 4096 /dev/sdc1
```

Create with explicit size:

```
dd if=/dev/zero of=test.hfs bs=1M count=50
mkfs.hfs -s 52428800 -L "Test" test.hfs
```

6.2 mkfs.hfs+ - HFS Plus Filesystem Creation

6.2.1 Synopsis

```
mkfs.hfs+ [-L label] [-l label] [-s size] [-b block_size]
          [-j] [-J journal_size] device
```

6.2.2 Description

Creates an HFS+ (Hierarchical File System Plus) volume. Supports journaling, larger volumes, and Unicode filenames. Initializes Volume Header, allocation bitmap, catalog B-tree, extents B-tree, attributes B-tree, and optionally journal.

6.2.3 Options

Option	Description
-L label	Set volume label (Unicode, primary option)
-l label	Set volume label (alias for -L)
-s size	Explicit volume size in bytes
-b block_size	Allocation block size (default: 4096, min 512)
-j	Enable journaling (WARNING: Linux kernel does not support)
-J size	Set journal size in bytes (requires -j)
device	Block device or file to format

Table 6.3: mkfs.hfs+ Options

6.2.4 Volume Label

Format: HFSUniStr255 (UTF-16BE, NFD normalized)

- Maximum: 255 UTF-16 characters (510 bytes + 2-byte length)

- Encoding: UTF-16 Big Endian

- Normalization: MUST be NFD (decomposed)

- Example: "Test" = 0x0004 0x0054 0x0065 0x0073 0x0074

6.2.5 Block Size Selection

Default: 4096 bytes (4 KB) - optimal for modern drives
Recommended values:

- 4 KB: Standard, best compatibility

- 8 KB: Large volumes (> 1 TB)

- 16 KB: Very large volumes (> 10 TB)

6.2.6 Initialization Sequence

Step 1: Device Validation

1. Open device read/write

2. Query device size

3. Verify minimum size (typically 1 MB)

4. Calculate total allocation blocks

5. Verify blocks fit in 32-bit field

Step 2: Volume Header Creation

1. Zero 512-byte Volume Header buffer

2. Set signature = 0x482B ('H+')

3. Set version = 4

4. Set attributes = 0x00000100 (unmounted bit)

5. Set lastMountedVersion = 0x6873786C ('hfsutil' signature)

6. Set journalInfoBlock = 0 (or block number if -j)

7. Set createDate, modifyDate = hfs_get_safe_time()

8. Set fileCount = 0, folderCount = 1 (root folder)

9. Set blockSize = selected block size

10. Set totalBlocks = calculated total

11. Set freeBlocks = total - reserved

12. Set rsrcClumpSize = blockSize * 4

13. Set dataClumpSize = blockSize * 4

14. Set nextCatalogID = 16

15. Initialize fork data for special files (allocation, extents, catalog, attributes)

Step 3: Allocation Bitmap File

1. Calculate bitmap size: (totalBlocks + 7) / 8 bytes

2. Allocate allocation file

3. Set logicalSize in Volume Header allocationFile

4. Mark boot blocks, VH, bitmap, B-trees as used

5. Write bitmap to allocation file extents

Step 4: Catalog B-Tree

1. Allocate catalog file (default: 20 allocation blocks)

2. Node size: 4096 bytes (standard)

3. Create header node (node 0):

   - BTHeaderRec: treeDepth = 1, nodeSize = 4096
   - btreeType = 128 (HFS+)
   - keyCompareType = 0xBC (case-insensitive)
   - maxKeyLength = 516 (2 + 4 + 510 for Unicode)

4. Create root folder record (CNID 2) as leaf node

5. Set fork data in volume Header catalogFile

Step 5: Extents B-Tree

1. Allocate extents file (default: 10 allocation blocks)

2. Create header node with empty tree

3. Initialize BTHeaderRec

4. No records initially (all files fit in catalog)

Step 6: Attributes B-Tree (Optional)

1. Normally left empty (logicalSize = 0)

2. Created on-demand when first attribute added

3. If created: allocate minimum space, initialize header

Step 7: Journal Setup (if -j)

1. Calculate journal size (default: 8 MB or 0.1% of volume)

2. Allocate journal file (CNID 14)

3. Create JournalInfoBlock at allocated block

4. Set journalInfoBlock in Volume Header

5. Set attributes —= 0x00002000 (journaled bit)

6. Initialize journal header with magic 0x4A4E4C78

7. **WARNING**: Emit warning about Linux incompatibility

Step 8: Alternate Volume Header

1. Copy Volume Header to buffer

2. Write to: device_size_bytes - 1024

3. Verify success

Step 9: Finalization

1. Sync all writes

2. Close device

3. Print summary (blocks, size, journal status)

4. Exit 0

6.2.7 Journaling Considerations

CRITICAL WARNING: The Linux HFS+ kernel driver does NOT support journaling.

- **macOS**: Full journal support, safe to use -j

- **Linux**: Journal ignored, potential corruption on crash

- **Recommendation**: Omit -j for Linux-mounted volumes

- **Alternative**: Disable journal on macOS before sharing with Linux

Journal size recommendations:

- Minimum: 1 MB

- Default: 8 MB or 0.1% of volume size (whichever larger)

- Maximum: 512 MB (practical limit)

- Formula: max(8MB, min(512MB, volume_size * 0.001))

6.2.8 Exit Codes

Code	Meaning
0	Success, filesystem created
1	Failure (invalid arguments, device error, insufficient space)

Table 6.4: mkfs.hfs+ Exit Codes (Unix Standard)

6.2.9 Examples

Create standard HFS+ volume:

```
mkfs.hfs+ -L "Data" /dev/sdb1
mkfs.hfs+ -L "Backup" backup.hfsplus
```

Create journaled volume (macOS only):

```
mkfs.hfs+ -j -L "Journal Test" /dev/sdc1
```

Create with custom block size:

```
mkfs.hfs+ -b 8192 -L "Large Volume" /dev/sdd1
```

Create image file:

```
dd if=/dev/zero of=test.hfsplus bs=1M count=100
mkfs.hfs+ -s 104857600 -L "Test Image" test.hfsplus
```

6.2.10 Verification After Creation

Always verify with fsck:

```
mkfs.hfs+ -L "New Volume" /dev/sdb1
fsck.hfs+ -n /dev/sdb1
```

Expected output: "Volume appears to be OK"

Check volume signature:

```
xxd -s 1024 -l 2 -p /dev/sdb1
# Expected: 482b (HFS+) or 4858 (HFSX)
```

6.2.11   Common Issues

Device Busy

Error: Device or resource busy

Solution: Unmount device first

```
umount /dev/sdb1
mkfs.hfs+ -L "Data" /dev/sdb1
```

Permission Denied

Error: Permission denied

Solution: Run with sudo

```
sudo mkfs.hfs+ -L "Data" /dev/sdb1
```

Volume Too Small

Error: Volume too small for HFS+

Solution: Use larger device or reduce block size

6.3   Implementation Details

6.3.1   Source Code Organization

- `src/mkfs/mkfs_hfs.c`: HFS Classic creation

- `src/mkfs/mkfs_hfsplus.c`: HFS+ creation

- `src/common/hfstime.c`: Time conversion utilities

- `src/common/version.c`: Version information

6.3.2   Critical Functions

hfs_get_safe_time():

```
uint32_t hfs_get_safe_time(void) {
    time_t now = time(NULL);
    uint32_t hfs_time = (uint32_t)now + 2082844800;
    if (hfs_time > 4294967295) {  // Y2K40
        hfs_time = 4102444800;     // Jan 1, 2030
    }
    return hfs_time;
}
```

calculate_block_size():

```
uint32_t calculate_block_size(uint64_t volume_size) {
    if (volume_size <= 256 * 1024 * 1024)
        return 512;
    if (volume_size <= 512 * 1024 * 1024)
        return 1024;
```

```
    if (volume_size <= 1024 * 1024 * 1024)
        return 2048;
    return 4096;  // Default for large volumes
}
```

6.4   Testing mkfs Utilities

See `test/test_mkfs.sh` for comprehensive test suite validating:

- Signature correctness (0x4244 for HFS, 0x482B for HFS+)

- Volume Header field initialization

- Allocation bitmap consistency

- B-tree initialization

- Alternate MDB/VH placement

- Exit code compliance

- fsck validation of created volumes

Chapter 7

fsck Utilities

Chapter 8

Filesystem Check and Repair Utilities

This chapter documents the fsck.hfs and fsck.hfs+ utilities for checking and repairing HFS and HFS+ filesystems.

8.1 fsck.hfs - HFS Classic Filesystem Check

8.1.1 Synopsis

```
fsck.hfs [-dfnpvy] [-b block_size] device
```

8.1.2 Description

Checks and optionally repairs an HFS (Hierarchical File System) Classic volume. Performs comprehensive validation of MDB, allocation bitmap, B-trees, and file/folder records. Can auto-detect HFS+ volumes and delegate to fsck.hfs+.

8.1.3 Options

Option	Description
-d	Debug mode (verbose output for developers)
-f	Force check even if volume appears clean
-n	No-modify mode (check only, do not repair)
-p	Preen mode (automatic safe repairs)
-v	Verbose output
-y	Yes to all repairs (assume yes to all prompts)
-b size	Override block size detection
device	Block device or file to check (required)

Table 8.1: fsck.hfs Options

8.1.4 Auto-Detection and Delegation

Critical feature: fsck.hfs automatically detects HFS+ volumes and delegates to fsck.hfs+
   Detection algorithm:

   1. Read bytes 1024-1025 (signature)

2. If signature == 0x482B ('H+') → Execute fsck.hfs+

3. If signature == 0x4858 ('HX') → Execute fsck.hfs+

4. If signature == 0x4244 ('BD') → Continue as HFS

5. Otherwise → Error: invalid filesystem

8.1.5 Validation Phases

Phase 1: MDB Validation

1. Read MDB at offset 1024

2. Verify drSigWord == 0x4244

3. Check drAtrb unmounted bit (0x0100)

4. Verify drNxtCNID $\geq$ 16

5. Validate drAlBlkSiz is power of 2

6. Check drNmAlBlks * drAlBlkSiz fits device

7. Verify drFreeBks $\leq$ drNmAlBlks

8. Check dates (createDate, modifyDate) for Y2K28

9. Validate volume name (drVN) is valid Pascal string

Phase 2: Alternate MDB Check

1. Calculate alternate MDB offset: device_size - 1024

2. Read alternate MDB

3. Compare with primary MDB

4. If mismatch: warn and offer repair

Phase 3: Allocation Bitmap Validation

1. Read allocation bitmap from drVBMSt

2. Verify bitmap size: (drNmAlBlks + 7) / 8 bytes

3. Count free blocks in bitmap

4. Compare with drFreeBks

5. If mismatch: error "Free block count incorrect"

6. Mark known allocated blocks:

   - Boot blocks (0-1)
   - MDB (block 2)
   - Bitmap itself
   - Catalog file extents
   - Extents file extents

Phase 4: Catalog B-Tree Validation

1. Read catalog file from extents in MDB

2. Validate header node (node 0):

   - Node descriptor: kind == 1
   - BTHeaderRec: validate all fields
   - bthDepth $\geq$ 1
   - bthRoot valid node number
   - bthNRecs matches actual count

3. Traverse B-tree:

   - Verify all node links (fLink, bLink)
   - Check key ordering (ascending)
   - Validate record types
   - Verify folder valence (file count)
   - Check CNID uniqueness

4. Validate root folder (CNID 2)

5. Check for orphaned files

Phase 5: Extents B-Tree Validation

1. Read extents file

2. Validate header node

3. Check all extent overflow records

4. Verify extents don't overlap

5. Ensure all extents within volume bounds

8.1.6   Repair Actions

Safe repairs (enabled with -p or -y):

- Fix free block count in bitmap

- Repair alternate MDB from primary

- Correct folder valence counts

- Fix minor B-tree inconsistencies

- Update modification date

  Unsafe repairs (require -y):

- Rebuild allocation bitmap

- Re-link orphaned files to lost+found

- Truncate files with invalid extents

- Rebuild B-tree structure

8.1.7   Exit Codes (Unix/BSD Standard)

Code	Meaning
0	Filesystem is clean, no errors
1	Filesystem errors corrected successfully
2	Filesystem errors corrected, reboot required
4	Filesystem errors left uncorrected
8	Operational error (cannot open device, etc.)
16	Usage error (invalid command line arguments)
32	Fsck canceled by user
128	Shared library error

Table 8.2: fsck.hfs Exit Codes (BSD Standard)

Note: These match BSD fsck standards for maximum compatibility.

8.1.8   Examples

Check without modifying:

```
fsck.hfs -n /dev/sdb1
```

Auto-repair in preen mode:

```
fsck.hfs -p /dev/sdb1
```

Force check and repair:

```
fsck.hfs -fy /dev/sdb1
```

Verbose check:

```
fsck.hfs -vn disk.hfs
```

8.2   fsck.hfs+ - HFS Plus Filesystem Check

8.2.1   Synopsis

```
fsck.hfs+ [-dfnpvy] [-b block_size] device
```

8.2.2   Description

Checks and repairs HFS+ (Hierarchical File System Plus) volumes. Handles journaling, Unicode filenames, larger volumes, and extended attributes. Performs journal replay if needed before checking.

8.2.3   Options

Same as fsck.hfs (see above table).

8.2.4   Journal Handling

Auto-detection:

1. Read Volume Header

2. Check attributes & 0x00002000 (journaled bit)

3. If journaled:

   - Read journalInfoBlock pointer
   - Load JournalInfoBlock
   - Read journal header
   - Verify magic == 0x4A4E4C78

Journal replay process:

1. Check journal dirty flag

2. If dirty:

   - Parse transaction log
   - Replay committed transactions
   - Write modified blocks to volume
   - Clear journal dirty flag
   - Emit message: "Journal replayed successfully"

3. Set Volume Header unmounted bit (0x0100)

Linux compatibility warning:

```
WARNING: Volume is journaled
Linux HFS+ driver does not support journaling
Mount read-only to prevent corruption
```

8.2.5   Validation Phases

Phase 1: Volume Header Validation

1. Read Volume Header at offset 1024

2. Verify signature == 0x482B or 0x4858

3. Check version == 4 or 5

4. Validate attributes field

5. Verify blockSize is power of 2

6. Check totalBlocks * blockSize $\leq$ device size

7. Verify freeBlocks $\leq$ totalBlocks

8. Check nextCatalogID $\geq$ 16

9. Validate rsrcClumpSize > 0

10. Validate dataClumpSize > 0

11. Check all fork data structures (allocation, extents, catalog, attributes)

12. Verify dates for Y2K40 (createDate, modifyDate, etc.)

Phase 2: Alternate Volume Header

1. Read from device_size - 1024

2. Compare with primary

3. If mismatch: offer repair

Phase 3: Allocation File Validation

1. Read allocation file from fork data

2. Verify bitmap size: (totalBlocks + 7) / 8

3. Count free blocks

4. Compare with freeBlocks in Volume Header

5. Mark known allocations:

    - Boot blocks
    - Volume Header
    - Allocation file itself
    - Extents file
    - Catalog file
    - Attributes file
    - Journal file (if journaled)

Phase 4: Catalog B-Tree Validation

1. Read catalog file from Volume Header

2. Validate header node:

    - BTHeaderRec complete validation
    - nodeSize == 4096 (standard)
    - btreeType == 128 (HFS+)
    - keyCompareType == 0xBC or 0xCF
    - maxKeyLength == 516

3. Traverse all nodes:

    - Validate node descriptor
    - Check record count vs actual
    - Verify key ordering (NFD Unicode)
    - Validate HFSUniStr255 strings
    - Check folder valence consistency
    - Verify file/folder CNIDs unique
    - Validate fork data in file records

4. Check root folder (CNID 2)

5. Verify hard links (if present)

6. Check Thread records consistency

Phase 5: Extents B-Tree Validation

1. Read extents file

2. Validate header

3. Check all overflow extent records

4. Verify no extent overlaps

5. Ensure extents within bounds

6. Cross-check with allocation bitmap

Phase 6: Attributes B-Tree (if present)

1. If attributesFile logicalSize > 0:

    - Read attributes file
    - Validate header
    - Check attribute records
    - Verify com.apple.* attribute validity

8.2.6   Repair Actions

Safe repairs (-p or -y):

- Fix free block count

- Repair alternate Volume Header

- Correct folder valence

- Fix minor B-tree issues

- Update Volume Header dates

- Clear bad unmounted flag

- Replay journal if dirty

    Aggressive repairs (-y only):

- Rebuild allocation bitmap

- Re-link orphaned files

- Rebuild catalog B-tree

- Fix extent overlaps

- Truncate corrupted files

8.2.7   Exit Codes

Same as fsck.hfs (BSD standard, see table above).

8.2.8   Examples

Check journaled volume:

```
fsck.hfs+ -n /dev/sdb1
# Output: Replaying journal...
# Output: Volume appears to be OK
```

Force check and repair:

```
fsck.hfs+ -fy /dev/sdb1
```

Preen mode (safe auto-repair):

```
fsck.hfs+ -p /dev/sdb1
```

Verbose check with debug:

```
fsck.hfs+ -dvn backup.hfsplus
```

8.2.9   Common Error Messages

Volume Needs Repair

```
** /dev/sdb1
   Volume check failed
   0 HFS Plus volume checked
```

Solution: Run with -y to repair

```
fsck.hfs+ -y /dev/sdb1
```

Journal Replay Failed

```
Journal replay failed: checksum mismatch
```

Solution: Journal corrupted, may need manual intervention. Try:

```
fsck.hfs+ -fy /dev/sdb1  # Force repair
```

Free Block Count Incorrect

```
Volume bitmap needs minor repair
Free block count is 12450, should be 12453
```

Solution: Safe to auto-repair with -p

Catalog B-Tree Corrupted

```
Catalog file entry not found for thread record
```

Solution: Serious error, requires -y repair

8.3   Implementation Details

8.3.1   Source Code Organization

- `src/fsck/fsck_hfs.c`: HFS Classic checker

- `src/fsck/fsck_hfsplus.c`: HFS+ checker

- `src/fsck/btree.c`: B-tree validation

- `src/fsck/journal.c`: Journal replay

- `src/common/hfstime.c`: Y2K40 safeguards

8.3.2   Critical Functions

check_volume_header():

```
int check_volume_header(struct hfs_vh *vh) {
    if (vh->signature != 0x482B && vh->signature != 0x4858)
        return -1;
    if (vh->version != 4 && vh->version != 5)
        return -1;
    if (vh->nextCatalogID < 16)
        return -1;
    if (vh->rsrcClumpSize == 0 || vh->dataClumpSize == 0)
        return -1;
    return 0;
}
```

replay_journal():

```
int replay_journal(int fd, struct hfs_vh *vh) {
    if (!(vh->attributes & 0x2000))
        return 0;  // Not journaled

    // Read JournalInfoBlock
    // Verify journal magic
    // Parse transactions
    // Write blocks to volume
    // Clear dirty flag

    return 0;  // Success
}
```

8.3.3   B-Tree Validation Algorithm

Recursive validation:

1. Start at root node (from BTHeaderRec)

2. For each node:

   - Validate node descriptor
   - Check record count
   - For index nodes: recurse to children

- For leaf nodes: validate records
- Verify key ordering

3. Track visited nodes (detect cycles)

4. Verify leaf chain (fLink/bLink)

8.4   Testing fsck Utilities

See `test/test_fsck.sh` for comprehensive tests:

- Create clean volume, verify exit code 0

- Create volume with errors, verify detection

- Test journal replay on dirty volume

- Verify all repair actions

- Test HFS+ delegation from fsck.hfs

- Validate exit code compliance

- Test -n (no modify) mode preserves data

8.4.1   Y2K40 Protection

Both fsck utilities use `hfs_get_safe_time()` to:

- Detect dates beyond Feb 6, 2040

- Emit warning if found

- Offer to correct to Jan 1, 2030

- Update Volume Header/MDB if repaired

Chapter 9

mount Utilities

Chapter 10

Mount Utilities

This chapter documents mount.hfs and mount.hfs+ utilities for mounting HFS and HFS+ filesystems following Unix/BSD/Linux standards.

10.1 mount.hfs and mount.hfs+ - Filesystem Mounting

10.1.1 Synopsis

```
mount.hfs [-o options] [-r] [-w] [-v] device mountpoint
mount.hfs+ [-o options] [-r] [-w] [-v] device mountpoint
mount -t hfs [-o options] device mountpoint
mount -t hfsplus [-o options] device mountpoint
```

10.1.2 Description

Mount HFS or HFS+ filesystems at specified mountpoints. These utilities follow Unix mount(8) conventions and integrate with /etc/fstab. Automatically validate filesystem before mounting and handle both read-only and read-write modes.

10.1.3 Standard Options

Option	Description
-o options	Mount options (comma-separated, see below)
-r	Mount read-only (shorthand for -o ro)
-w	Mount read-write (shorthand for -o rw)
-v	Verbose output
device	Block device to mount (e.g., /dev/sdb1)
mountpoint	Directory where filesystem will be mounted

Table 10.1: mount.hfs/mount.hfs+ Options

10.1.4 Mount Options (-o)

Option	Description
ro	Read-only mount (default for journaled on Linux)
rw	Read-write mount
noexec	Do not allow execution of binaries

Option	Description
nosuid	Do not honor setuid/setgid bits
nodev	Do not interpret device files
sync	Synchronous I/O (slow but safe)
async	Asynchronous I/O (fast but risky on crash)
uid=N	Set owner UID for all files
gid=N	Set group GID for all files
umask=N	Set umask for file permissions
force	Force mount even if filesystem appears dirty

Table 10.2: Common Mount Options

10.1.5   Mount Process

Pre-Mount Validation

1. Verify device exists and is accessible

2. Check if device is already mounted

3. Read filesystem signature (offset 1024)

4. Validate signature (0x4244 for HFS, 0x482B/0x4858 for HFS+)

5. Check unmounted bit in attributes

6. If dirty: warn and suggest fsck

7. Verify mountpoint exists and is empty directory

Kernel Module Check (Linux)

```
# Check if kernel module loaded
lsmod | grep hfs
# Load if needed
modprobe hfs       # For HFS classic
modprobe hfsplus  # For HFS+
```

Mount Execution

1. Call kernel mount() syscall

2. Pass filesystem type ("hfs" or "hfsplus")

3. Pass mount options

4. Kernel driver validates and mounts

5. Update /etc/mtab (on success)

6. Return exit code

10.1.6   Journaling Considerations

CRITICAL for Linux users:

- Linux HFS+ driver does NOT support journaling

- Journaled volumes MUST be mounted read-only on Linux

- Automatic detection: if journaled bit set $\rightarrow$ force ro

- Warning displayed: "Volume is journaled, mounting read-only"

Detection algorithm:

```
if (attributes & 0x00002000) {
    fprintf(stderr, "WARNING: Journaled volume\n");
    fprintf(stderr, "Linux does not support HFS+ journaling\n");
    fprintf(stderr, "Mounting read-only\n");
    force_readonly = 1;
}
```

10.1.7   Exit Codes

Code	Meaning
0	Success, filesystem mounted
1	Incorrect invocation or permissions
2	System error (out of memory, etc.)
4	Internal mount bug
8	User interrupt
16	Problems writing or locking /etc/mtab
32	Mount failure
64	Some mount succeeded

Table 10.3: mount Exit Codes (Linux Standard)

10.1.8   Examples

Mount HFS classic read-write:

```
mount.hfs /dev/sdb1 /mnt/hfs
mount -t hfs /dev/sdb1 /mnt/hfs
```

Mount HFS+ read-only:

```
mount.hfs+ -r /dev/sdc1 /mnt/backup
mount -t hfsplus -o ro /dev/sdc1 /mnt/backup
```

Mount with specific permissions:

```
mount.hfs+ -o uid=1000,gid=1000,umask=022 /dev/sdd1 /mnt/data
```

Force mount dirty volume (WARNING - DANGEROUS):

```
mount.hfs+ -o force /dev/sde1 /mnt/recovery
# WARNING: Can cause data corruption!
```

10.1.9   Unmounting

Standard unmount:

```
umount /mnt/hfs
umount /dev/sdb1
```

Force unmount (if busy):

```
umount -f /mnt/hfs
# Or lazy unmount:
umount -l /mnt/hfs
```

Check what's using mount:

```
lsof | grep /mnt/hfs
fuser -m /mnt/hfs
```

10.1.10   /etc/fstab Integration

HFS entry:

```
/dev/sdb1  /mnt/hfs  hfs  defaults,ro  0  0
```

HFS+ entry:

```
/dev/sdc1  /mnt/backup  hfsplus  ro,noexec,nosuid  0  0
```

HFS+ with UUID:

```
UUID=xxxx-xxxx  /mnt/data  hfsplus  rw,uid=1000  0  2
```

Field meanings:

1. Device or UUID

2. Mount point

3. Filesystem type (hfs or hfsplus)

4. Mount options

5. Dump frequency (usually 0)

6. fsck pass number (0=skip, 1=root, 2=other)

10.1.11   Common Issues

Device is Busy

```
mount: /dev/sdb1 already mounted or /mnt/hfs busy
```

Solutions:

- Check if already mounted: `mount | grep sdb1`

- Find processes using it: `lsof | grep sdb1`

- Kill processes or use `umount -f`

Wrong Filesystem Type

`mount: wrong fs type, bad option, bad superblock`

Solutions:

- Verify signature: `xxd -s 1024 -l 2 /dev/sdb1`

- Use correct type: hfs vs hfsplus

- Run fsck first: `fsck.hfs+ /dev/sdb1`

Permission Denied

`mount: only root can do that`

Solution: Use sudo

`sudo mount.hfs+ /dev/sdb1 /mnt/hfs`

Module Not Found

`mount: unknown filesystem type 'hfsplus'`

Solution: Load kernel module

```
sudo modprobe hfsplus
# Or for HFS classic:
sudo modprobe hfs
```

10.1.12   Portability Notes

Linux: Full support via kernel modules (hfs.ko, hfsplus.ko)

- HFS: Full read/write support

- HFS+: Read/write support, NO journaling

- Always mount journaled volumes read-only

  FreeBSD: Native HFS read support

- Write support via FUSE

- Mount with: `mount -t hfs /dev/da0s1 /mnt`

  macOS: Native full support

- Automatic mounting via diskutil

- Full journaling support

- Default filesystem (legacy, now APFS)

10.2   Implementation

10.2.1   Source Code

- `src/mount/mount_hfs.c`: HFS mount helper

- `src/mount/mount_hfsplus.c`: HFS+ mount helper

- `src/mount/mount_common.h`: Shared definitions

10.2.2   Critical Functions

validate_filesystem():

```
int validate_filesystem(const char *device) {
    int fd = open(device, O_RDONLY);
    uint16_t sig;
    lseek(fd, 1024, SEEK_SET);
    read(fd, &sig, 2);
    sig = be16toh(sig);

    if (sig == 0x4244) return FS_HFS;
    if (sig == 0x482B || sig == 0x4858) return FS_HFSPLUS;
    return FS_UNKNOWN;
}
```

check_journaling():

```
int check_journaling(int fd) {
    uint32_t attributes;
    lseek(fd, 1024 + 4, SEEK_SET);
    read(fd, &attributes, 4);
    attributes = be32toh(attributes);
    return (attributes & 0x00002000) ? 1 : 0;
}
```

Chapter 11

hfsutil Commands

Chapter 12

hfsutil Commands

This chapter documents the hfsutil suite of commands for interactive manipulation of HFS/HFS+ volumes without mounting.

12.1 Overview

hfsutil provides a set of Unix-like commands for working with HFS and HFS+ volumes directly, similar to mtools for FAT filesystems. All commands operate on unmounted volumes.

12.1.1 Common Features

- Work on HFS and HFS+ volumes

- No mounting required

- Unix-style command interface

- Support for both block devices and image files

- MacRoman to UTF-8 conversion

12.2 Volume Management Commands

12.2.1 hformat - Format HFS Volume

Synopsis:

```
hformat [-l label] device
```

> **Description**: Creates a new HFS Classic volume (equivalent to mkfs.hfs).
> **Options**:

- -l label: Set volume label (1-27 characters)

- device: Block device or file to format

Example:

```
hformat -l "BackupDisk" /dev/sdb1
hformat -l "Test" disk.hfs
```

12.2.2 hmount - Mount HFS Volume (hfsutil context)

Synopsis:

```
hmount device [partition]
```

Description: Register HFS volume for use with other hfsutil commands (does NOT mount to OS).

Options:

- device: Block device or image file

- partition: Optional partition number (for APM)

Example:

```
hmount /dev/sdb1
hmount disk.hfs
hmount /dev/sdc 2  # Partition 2
```

12.2.3 humount - Unmount HFS Volume (hfsutil context)

Synopsis:

```
humount [device]
```

Description: Unregister HFS volume from hfsutil.

Example:

```
humount
humount /dev/sdb1
```

12.2.4 hvol - Display Volume Information

Synopsis:

```
hvol [device]
```

Description: Show volume information (name, creation date, free space, etc.).

Example:

```
hvol
# Output:
# Volume name: MyDisk
# Volume created: Mon Jan 15 12:00:00 2024
# Total blocks: 10240
# Free blocks: 5120
```

12.3 File Operations Commands

12.3.1 hls - List Directory

Synopsis:

```
hls [-1abcdefgilmnrstux] [path ...]
```

Description: List files and directories (like Unix ls).

Common Options:

- -l: Long format (permissions, size, date)

- -a: Show all files (including invisible)

- -i: Show file IDs (CNIDs)

- -d: List directory itself, not contents

- -R: Recursive listing

- -1: One file per line

Examples:

```
hls                   # List current directory
hls -l                # Long format
hls -la /             # All files in root, long format
hls -R /System        # Recursive list
```

12.3.2   hcopy - Copy Files

Synopsis:

```
hcopy [-m|-b|-t|-r] source-path target-path
hcopy [-m|-b|-t|-r] source-path [...] target-directory
```

Description: Copy files to/from HFS volume.
Fork Options:

- -m: MacBinary format (default, includes both forks)

- -b: Copy both data and resource forks

- -t: Text mode (data fork only, line ending conversion)

- -r: Raw mode (data fork only, no conversion)

Examples:

```
# Copy from HFS to Unix:
hcopy :file.txt ./file.txt
hcopy -r :document.doc ./document.doc

# Copy to HFS:
hcopy ./readme.txt :README.TXT
hcopy ./file1 ./file2 :Folder/

# MacBinary format (preserves resource fork):
hcopy -m :app.bin ./app.bin
```

12.3.3   hmkdir - Create Directory

Synopsis:

```
hmkdir path [...]
```

Description: Create new directories on HFS volume.
Example:

```
hmkdir :Documents
hmkdir :Files :Backup :Archives
```

12.3.4 hdel - Delete Files/Directories

Synopsis:

```
hdel path [...]
```

    Description: Delete files and directories (directories must be empty).
    Example:

```
hdel :oldfile.txt
hdel :TempFolder
hdel :file1.doc :file2.doc
```

12.3.5 hrename - Rename Files

Synopsis:

```
hrename src-path dest-path
```

    Description: Rename or move files/directories within volume.
    Example:

```
hrename :oldname.txt :newname.txt
hrename :File.doc :Archive/File.doc
```

12.4 Navigation Commands

12.4.1 hcd - Change Directory

Synopsis:

```
hcd [path]
```

    Description: Change current directory in HFS volume.
    Example:

```
hcd :Documents
hcd ..
hcd /
```

12.4.2 hpwd - Print Working Directory

Synopsis:

```
hpwd
```

    Description: Display current directory path.
    Example:

```
hp
```

```
wd
# Output: :Documents:Projects
```

12.5   Attribute Commands

12.5.1   hattrib - Show/Modify Attributes

Synopsis:

```
hattrib [-t TYPE] [-c CREA] [-i|+i] path [...]
```

Description: Display or modify file/folder attributes (type, creator, invisible flag).
Options:

- -t TYPE: Set file type (4 characters)

- -c CREA: Set creator (4 characters)

- -i: Make invisible

- +i: Make visible

Examples:

```
hattrib :file.txt              # Show attributes
hattrib -t TEXT -c EDIT :file.txt  # Set type/creator
hattrib -i :HiddenFile         # Make invisible
hattrib +i :NowVisible          # Make visible
```

12.6   Path Syntax

HFS path format:

- : prefix indicates HFS volume path

- : separator between folders (not /)

- Example: :Folder:Subfolder:file.txt

- / is NOT a path separator in HFS

- Root: : or :/

Examples:

```
:                   # Root directory
:Documents          # Documents folder in root
:Docs:file.txt      # file.txt in Docs folder
..                  # Parent directory
```

12.7   Character Encoding

HFS uses MacRoman encoding:

- hfsutil converts MacRoman ↔ UTF-8 automatically

- Some characters may not convert perfectly

- Colon (:) forbidden in filenames (path separator)

12.8   Workflow Example

Complete workflow for accessing HFS volume:

```
# 1. Mount volume to hfsutil
hmount /dev/sdb1

# 2. Navigate and explore
hpwd
hls -la
hcd :Documents

# 3. Copy files
hcopy :file.txt ./backup/
hcopy ./newfile.doc :

# 4. Create directory and organize
hmkdir :Archive
hrename :oldfile.txt :Archive/oldfile.txt

# 5. Cleanup
hdel :tempfile.tmp

# 6. Unmount
humount
```

12.9   Testing

See `test/test_hfsutils.sh` for comprehensive tests:

- Format, mount, unmount sequence

- File creation and deletion

- Directory operations

- Copy to/from volume

- Attribute manipulation

- Path handling

Chapter 13

Implementation Details

Chapter 14

Implementation Details

14.1 Source Code Organization

14.1.1 Directory Structure

```
hfsutils/
 src/
    common/           # Shared utilities
       hfstime.c    # Time conversion
       endian.h     # Endian handling
       version.c    # Version info
    mkfs/             # Filesystem creation
       mkfs_hfs.c   # HFS creator
       mkfs_hfsplus.c # HFS+ creator
    fsck/             # Filesystem checker
       fsck_hfs.c
       fsck_hfsplus.c
       btree.c      # B-tree validation
       journal.c    # Journal replay
    mount/            # Mount helpers
       mount_hfs.c
       mount_hfsplus.c
    hfsutil/          # HFS utilities
        hformat.c
        hmount.c
        hcopy.c
 test/                # Test suite
    test_mkfs.sh
    test_fsck.sh
    test_hfsutils.sh
 doc/                 # Documentation
     man/
      latex/
```

14.2 Critical Algorithms

14.2.1 B-Tree Traversal

```
void traverse_btree(BTNodePtr node, int depth) {
    if (node->kind == kBTLeafNode) {
```

```
        validate_leaf_records(node);
        return;
    }

    for (int i = 0; i < node->numRecords; i++) {
        BTNodePtr child = read_child_node(node, i);
        traverse_btree(child, depth + 1);
    }
}
```

14.2.2   Extent Allocation

```
int allocate_extents(VolumePtr vol, uint32_t blocks_needed,
                     HFSPlusExtentRecord extents) {
    uint32_t blocks_found = 0;
    uint32_t start_block = vol->nextAllocation;

    while (blocks_found < blocks_needed) {
        uint32_t contig = find_contiguous_blocks(vol, start_block);
        if (contig == 0) return -1; // No space

        add_extent(extents, start_block, contig);
        blocks_found += contig;
        start_block += contig;
    }

    return 0;
}
```

14.2.3   Unicode NFD Normalization

```
void normalize_to_nfd(uint16_t *unicode, size_t *length) {
    for (size_t i = 0; i < *length; i++) {
        uint16_t ch = unicode[i];
        const DecompEntry *decomp = lookup_decomp(ch);

        if (decomp) {
            // Replace with decomposed form
            memmove(&unicode[i+decomp->count], &unicode[i+1],
                    (*length - i - 1) * sizeof(uint16_t));
            memcpy(&unicode[i], decomp->chars,
                    decomp->count * sizeof(uint16_t));
            *length += (decomp->count - 1);
            i += (decomp->count - 1);
        }
    }
}
```

14.3   Data Structure Handling

14.3.1   Endianness Conversion

```
#define be16toh(x) ntohs(x)
#define be32toh(x) ntohl(x)
#define htobe16(x) htons(x)
#define htobe32(x) htonl(x)

void read_volume_header(int fd, HFSPlusVolumeHeader *vh) {
    lseek(fd, 1024, SEEK_SET);
    read(fd, vh, sizeof(*vh));

    // Convert all multi-byte fields
    vh->signature = be16toh(vh->signature);
    vh->version = be16toh(vh->version);
    vh->attributes = be32toh(vh->attributes);
    vh->blockSize = be32toh(vh->blockSize);
    vh->totalBlocks = be32toh(vh->totalBlocks);
    // ... convert all other fields
}
```

14.3.2   Memory Safety

- All buffers bounds-checked

- No use-after-free (valgrind verified)

- No memory leaks in normal operation

- Defensive programming throughout

14.4   Performance Optimizations

- Bitmap operations use bitwise ops

- B-tree nodes cached during traversal

- Extent descriptors packed efficiently

- Minimal syscalls (buffered I/O)

Chapter 15

Testing and Validation

Chapter 16

Testing and Validation

16.1 Zero-Tolerance Testing Philosophy

Core principle: ALL filesystems must be 100% specification-compliant. ANY deviation results in test failure.

16.1.1 Why Zero Tolerance

- Filesystems store critical user data

- Even minor corruption can cause data loss

- Spec compliance ensures interoperability

- No "good enough" - either correct or broken

16.2 Test Suite Organization

16.2.1 test_mkfs.sh - Filesystem Creation Tests

Coverage:

- HFS signature validation (0x4244)

- HFS+ signature validation (0x482B)

- MDB field initialization (all 162 bytes)

- Volume Header initialization (all 512 bytes)

- Allocation block calculations

- B-tree header node creation

- Alternate MDB/VH placement

- Exit code compliance (0 = success, 1 = failure)

  Validation method:

```
# Create filesystem
mkfs.hfs+ -L "Test" test.img

# Verify signature
```

```
SIG=$(xxd -s 1024 -l 2 -p test.img)
if [ "$SIG" != "482b" ]; then
    echo "FAIL: Invalid signature"
    exit 1
fi

# Run fsck validation
fsck.hfs+ -n test.img
if [ $? -ne 0 ]; then
    echo "FAIL: fsck detected errors"
    exit 1
fi
```

16.2.2   test_fsck.sh - Validation Tests

Coverage:

- Clean volume detection (exit 0)

- Dirty volume detection

- MDB/Volume Header validation

- Allocation bitmap consistency

- B-tree structure validation

- Catalog integrity checks

- Journal detection and replay

- Exit code compliance (0, 1, 2, 4, 8)

  Test scenarios:

  1. Create clean volume → fsck returns 0

  2. Modify free block count → fsck detects error

  3. Create journaled volume → fsck replays journal

  4. Test -n (no modify) mode

  5. Verify repair actions with -y

16.2.3   test_hfsutils.sh - Command Tests

Coverage:

- hformat volume creation

- hmount/humount sequence

- hls directory listing

- hcopy file transfer

- hmkdir directory creation

- hdel file deletion

- Path handling (: separator)

- MacRoman encoding conversion

Example test:

```
# Format volume
hformat -l "TestVol" test.hfs

# Mount
hmount test.hfs

# Create directory
hmkdir :TestDir

# Copy file
echo "test" > temp.txt
hcopy temp.txt :TestDir/file.txt

# Verify
hls -l :TestDir | grep file.txt
if [ $? -ne 0 ]; then
    echo "FAIL: File not found"
    exit 1
fi

# Cleanup
humount
```

16.3   Continuous Integration

Automated testing on:

- Every commit

- Pull requests

- Multiple platforms (Linux, FreeBSD)

- Multiple architectures (x86_64, ARM)

16.4   Regression Prevention

All bugs get tests:

- Bug discovered → Test added

- Test fails until bug fixed

- Test remains in suite forever

- Prevents regression

16.5   Manual Testing Checklist

Before release:

1. All automated tests pass

2. Create HFS volume, mount on macOS

3. Create HFS+ volume, mount on Linux

4. Test large files (> 1 GB)

5. Test many files ($> 10{,}000$)

6. Test long filenames (255 chars)

7. Test special characters

8. Test journal replay

9. Verify fsck repairs

10. Cross-platform compatibility

Appendix A

Appendix

Appendix B

Appendix

B.1 C Structure Definitions

B.1.1 HFS Master Directory Block

```
struct HFSMasterDirectoryBlock {
    uint16_t  drSigWord;       // 0x4244
    uint32_t  drCrDate;        // Creation date
    uint32_t  drLsMod;         // Last mod date
    uint16_t  drAtrb;          // Attributes
    uint16_t  drNmFls;         // File count
    uint16_t  drVBMSt;         // Bitmap start block
    uint16_t  drAllocPtr;      // Alloc search start
    uint16_t  drNmAlBlks;      // Total alloc blocks
    uint32_t  drAlBlkSiz;      // Alloc block size
    uint32_t  drClpSiz;        // Clump size
    uint16_t  drAlBlSt;        // First alloc block
    uint32_t  drNxtCNID;       // Next CNID
    uint16_t  drFreeBks;       // Free blocks
    uint8_t   drVN[28];        // Volume name (Pascal)
    // ... (total 162 bytes)
} __attribute__((packed));
```

B.1.2 HFS+ Volume Header

```
struct HFSPlusVolumeHeader {
    uint16_t  signature;           // 0x482B or 0x4858
    uint16_t  version;             // 4 or 5
    uint32_t  attributes;          // Volume attributes
    uint32_t  lastMountedVersion;  // OS signature
    uint32_t  journalInfoBlock;    // Journal block (or 0)
    uint32_t  createDate;          // HFS+ time
    uint32_t  modifyDate;
    uint32_t  backupDate;
    uint32_t  checkedDate;
    uint32_t  fileCount;
    uint32_t  folderCount;
    uint32_t  blockSize;           // Bytes
    uint32_t  totalBlocks;
    uint32_t  freeBlocks;
```

```
    uint32_t  nextAllocation;
    uint32_t  rsrcClumpSize;
    uint32_t  dataClumpSize;
    uint32_t  nextCatalogID;        // >= 16
    uint32_t  writeCount;
    uint64_t  encodingsBitmap;
    uint32_t  finderInfo[8];
    HFSPlusForkData allocationFile;  // 80 bytes each
    HFSPlusForkData extentsFile;
    HFSPlusForkData catalogFile;
    HFSPlusForkData attributesFile;
    HFSPlusForkData startupFile;
} __attribute__((packed));  // 512 bytes total
```

B.1.3   HFS+ Fork Data

```
struct HFSPlusForkData {
    uint64_t  logicalSize;          // File size in bytes
    uint32_t  clumpSize;
    uint32_t  totalBlocks;
    HFSPlusExtentDescriptor extents[8];  // 8 x 8 bytes
} __attribute__((packed));  // 80 bytes total

struct HFSPlusExtentDescriptor {
    uint32_t  startBlock;
    uint32_t  blockCount;
} __attribute__((packed));  // 8 bytes
```

B.1.4   B-Tree Node Descriptor

```
struct BTNodeDescriptor {
    uint32_t  fLink;        // Forward link
    uint32_t  bLink;        // Backward link
    int8_t    kind;         // -1=leaf, 0=index, 1=header, 2=map
    uint8_t   height;       // 0 for leaves
    uint16_t  numRecords;
    uint16_t  reserved;
} __attribute__((packed));  // 14 bytes
```

B.1.5   B-Tree Header Record

```
struct BTHeaderRec {
    uint16_t  treeDepth;
    uint32_t  rootNode;
    uint32_t  leafRecords;
    uint32_t  firstLeafNode;
    uint32_t  lastLeafNode;
    uint16_t  nodeSize;         // Usually 4096
    uint16_t  maxKeyLength;
    uint32_t  totalNodes;
    uint32_t  freeNodes;
    uint16_t  reserved1;
```

```
    uint32_t  clumpSize;
    uint8_t   btreeType;        // 0=HFS, 128=HFS+
    uint8_t   keyCompareType;   // 0xBC or 0xCF
    uint32_t  attributes;
    uint32_t  reserved3[16];
} __attribute__((packed));  // 106 bytes
```

B.1.6   HFS+ Catalog File Record

```
struct HFSPlusCatalogFile {
    int16_t   recordType;       // 0x0002
    uint16_t  flags;
    uint32_t  reserved1;
    uint32_t  fileID;           // CNID
    uint32_t  createDate;
    uint32_t  contentModDate;
    uint32_t  attributeModDate;
    uint32_t  accessDate;
    uint32_t  backupDate;
    HFSPlusBSDInfo permissions;
    FInfo     userInfo;
    FXInfo    finderInfo;
    uint32_t  textEncoding;
    uint32_t  reserved2;
    HFSPlusForkData dataFork;
    HFSPlusForkData resourceFork;
} __attribute__((packed));  // 248 bytes
```

B.2   Error Codes

B.2.1   mkfs Exit Codes

Code	Meaning
0	Success
1	Failure (any error)

Table B.1: mkfs.hfs/mkfs.hfs+ Exit Codes

B.2.2   fsck Exit Codes

Code	Meaning
0	No errors
1	Errors corrected
2	Errors corrected, reboot needed
4	Errors uncorrected
8	Operational error
16	Usage error
32	Canceled
128	Library error

Code	Meaning

Table B.2: fsck.hfs/fsck.hfs+ Exit Codes (BSD Standard)

B.2.3 mount Exit Codes

Code	Meaning
0	Success
1	Incorrect invocation
2	System error
32	Mount failure

Table B.3: mount.hfs/mount.hfs+ Exit Codes

B.3 Glossary

Allocation Block Fundamental unit of disk space allocation in HFS/HFS+. Size ranges from 512 bytes to 64 KB.

Alternate MDB/VH Backup copy of Master Directory Block or Volume Header located at end of volume (offset: size - 1024 bytes).

B-tree Balanced tree data structure used for catalog, extents, and attributes files.

Big-Endian Byte order where most significant byte comes first. Used by HFS/HFS+ (Motorola format).

Catalog B-tree containing all files and folders on volume. CNID 4 in HFS+.

CNID Catalog Node ID - unique identifier for each file/folder. Reserved: 1-15, first user: 16.

Extent Contiguous range of allocation blocks. Described by startBlock + blockCount.

Fork Part of a file. HFS+ supports data fork and resource fork (80 bytes each in catalog).

HFSUniStr255 Unicode string format: 2-byte length + up to 255 UTF-16BE characters. MUST be NFD normalized.

Journal Transaction log for crash recovery. Optional in HFS+. NOT supported by Linux kernel driver.

MDB Master Directory Block - main metadata structure for HFS Classic. 162 bytes at offset 1024.

NFD Normalization Form D (Decomposed) - required Unicode form for HFS+ filenames. Example: é = e + combining acute.

Pascal String Length-prefixed string (1 byte length + characters). Used in HFS for volume names.

Volume Header Main metadata structure for HFS+. 512 bytes at offset 1024. Contains all filesystem parameters.

Y2K28 HFS Classic date overflow on February 6, 2028 (32-bit seconds since 1904).

Y2K40 HFS+ date overflow on February 6, 2040 (32-bit seconds since 1904).

B.4 References

B.4.1 Primary Sources

1. **Apple Technical Note TN1150** - "HFS Plus Volume Format"

2. **Inside Macintosh: Files** - Original HFS specification

3. **Linux Kernel Source** - fs/hfs/ and fs/hfsplus/ drivers

4. **FreeBSD Source** - sys/fs/hfs/

B.4.2 This Documentation

This manual provides complete bit-level specifications enabling reimplementation without external references. Required external resources:

- Unicode NFD decomposition tables (standard Unicode data)

- B-tree algorithms (standard CS textbook)

- CRC32 implementation (standard algorithm)

  Internet NOT required for implementation after obtaining these standard resources.

B.4.3 Online Resources (Optional)

- https://developer.apple.com/legacy/library/technotes/tn/tn1150.html

- Linux kernel documentation: Documentation/filesystems/hfsplus.txt

- Unicode tables: https://unicode.org/Public/UNIDATA/

B.5 Version History

hfsutils 4.1.0A.2:

- Complete LaTeX documentation (this manual)

- Chapters 00-10: 5,000 lines of bit-level specifications

- mkfs, fsck, mount, hfsutil fully documented

- Zero internet dependency for reimplementation